

# A Survey of Performance Optimization for Mobile Applications

Max Hort<sup>ID</sup>, Maria Kechagia<sup>ID</sup>, Federica Sarro<sup>ID</sup>, and Mark Harman

**Abstract**—To ensure user satisfaction and success of mobile applications, it is important to provide highly performant applications. This is particularly important for resource-constrained systems such as mobile devices. Thereby, non-functional performance characteristics, such as energy and memory consumption, play an important role for user satisfaction. This paper provides a comprehensive survey of non-functional performance optimization for Android applications. We collected 156 unique publications, published between 2008 and 2020, that focus on the optimization of performance of mobile applications. We target our search at four performance characteristics: responsiveness, launch time, memory and energy consumption. For each performance characteristic, we categorize optimization approaches based on the method used in the corresponding publications. Furthermore, we identify research gaps in the literature for future work.

**Index Terms**—Mobile applications, Android, non-functional performance optimization, software optimization, literature survey

## 1 INTRODUCTION

THE relevance of mobile (handheld) devices, such as the so-called *smartphones*, has been ever growing for the past ten years, reaching an estimate of 3.2 billion smartphone users in 2019.<sup>1</sup> Smartphones can nowadays be considered as the main information processing devices for users. With smartphones, users cannot only receive and make phone calls, but also execute similar tasks as those performed on personal computers (e.g., surf the internet, perform calculations, pay bills).

Even though mobile devices are powerful, they represent resource-constrained devices making the development of applications that can run on them (*mobile applications*) challenging. This means that the functionality and performance of mobile applications depend on the characteristics of mobile phones (e.g., their physical memory, processors, battery) and on the current execution context (e.g., how many applications run at the same time on a mobile phone).

To ensure the success of an application (e.g., whether it will be used, updated, or uninstalled [1], [2]), developers aim to maximize user experience quality, and, consequently, user satisfaction [3], [4], [5]. User satisfaction is mainly influenced by functional (*Does the application operate*

*as the user expects?*) and non-functional (*How does the application perform?*) application characteristics [1], [2]. Examples of functional issues can include missing or buggy features (e.g., a game application that functions in a different way than presented in its description). An example of a non-functional characteristic is the energy consumption of an application. Regardless of an application's functionality, users will be dissatisfied if the application drains the battery of their mobile devices within minutes.

Finkelstein *et al.* [6] found that the success of mobile applications in terms of downloads is correlated to the rating that the application attracts. These ratings are recorded by App Stores (e.g., Google Play, Apple Store, BlackBerry World). In 2018, the number of total applications downloaded amounted to 194 billion,<sup>2</sup> with every user having a multitude of different applications installed on their phone [7], [8]. With such a high number of applications, 75 percent of mobile device usage is filled by mobile applications [5]. While several studies show the importance of fixing software bugs that hinder applications' smooth function [9], [10], [11], [12], non-functional performance characteristics have shown to have a strong impact on user satisfaction as well [1], [2], [3], [13], [14], [15], [16], [17], [18]. This impact can be seen in the user reviews of real-world mobile applications:

- “This app is destroying my battery. I will have to uninstall it if there isn't a fix soon.” [13]
- “It lags and doesn't respond to my touch which almost always causes me to run into stuff.” [19]
- “Bring back the old version. Scrolling lags.” [3]
- “Makes GPS stay on all the time. Kills my battery.” [3]
- “Too much memory usage for a glorified web portal ad machine.” [18]

1. <https://www.statista.com/statistics/330695/number-of-smartphone-users-worldwide/>

• Max Hort, Maria Kechagia, and Federica Sarro are with the Department of Computer Science, University College London, WC1E 6BT London, U.K. E-mail: {max.hort.19, m.kechagia, f.sarro}@ucl.ac.uk.  
• Mark Harman is with the Department of Computer Science, University College London, WC1E 6BT London, U.K., and also with Facebook London, W1T 1FB, London, U.K. E-mail: mark.harman@ucl.ac.uk.

Manuscript received 13 June 2020; revised 18 Mar. 2021; accepted 28 Mar. 2021.  
Date of publication 6 Apr. 2021; date of current version 15 Aug. 2022.  
(Corresponding author: Federica Sarro.)

Recommended for acceptance by N. Nagappan.

Digital Object Identifier no. 10.1109/TSE.2021.3071193

2. <https://www.statista.com/statistics/271644/worldwide-free-and-paid-mobile-app-store-downloads/>

Furthermore, Banerjee and Roychoudhury [20] conducted a study on 170,000 user reviews, and showed that poor performance and energy consumption lead to application downvotes from users. Among all causes of downvotes, energy consumption caused the highest ratio of uninstallations.

Given the importance of non-functional performance characteristics on user satisfaction and the consequent success of mobile applications, as well as new optimization approaches that are developed each year, we provide a comprehensive overview of existing approaches for non-functional performance optimization of mobile applications. This can be used by practitioners, developers, and researchers to search for approaches appropriate to their needs (e.g., “How can I reduce energy consumption only by changing application source code?” or “Can I improve responsiveness by applying changes to the device hardware?”). Furthermore, we provide information on dependencies among non-functional properties, which reside in mobile applications. We focus our review on the Android platform, since it is open-source software and has the highest market share among mobile platforms at the time of writing this survey.<sup>3</sup>

Initially, we gathered and analyzed existing work to detect non-functional performance characteristics (Section 3). Based on this, we identified four non-functional characteristics, which describe user-perceived *performance* of mobile applications, and thereby applications success, i.e.: *responsiveness*, *launch time*, *memory consumption*, and *energy consumption*. For each of these characteristics, we have categorized previous work based on the optimization level (e.g., optimization applied to application, platform, or hardware level) and proposed optimization type (e.g., prefetching, preloading, display), as shown in Fig. 1.

We found that the majority of approaches to optimize responsiveness applied changes to the application’s source code, while launch time was improvement by changes to the Android platform. Approaches that optimize memory apply changes to both the application and Android platform’s source code. The majority of work was concerned with optimizing energy consumption. Moreover, we were able to detect relationships among the four non-functional performance characteristics (e.g., energy consumption can increase with an improved responsiveness of applications).

To the best of our knowledge, this is the first survey to investigate multiple non-functional performance characteristics and their relationships. To summarize, our work:

- (1) provides a comprehensive literature review of the state-of-the-art research on the optimization of non-functional characteristics for mobile applications;
- (2) provides a categorization of existing optimization approaches based on their level and type;
- (3) identifies challenges and opportunities for future research in this area.

We have made publicly available additional resources [21] and an online version of the work reviewed in this survey, which we aim to keep up-to-date by accepting external contributions [21].

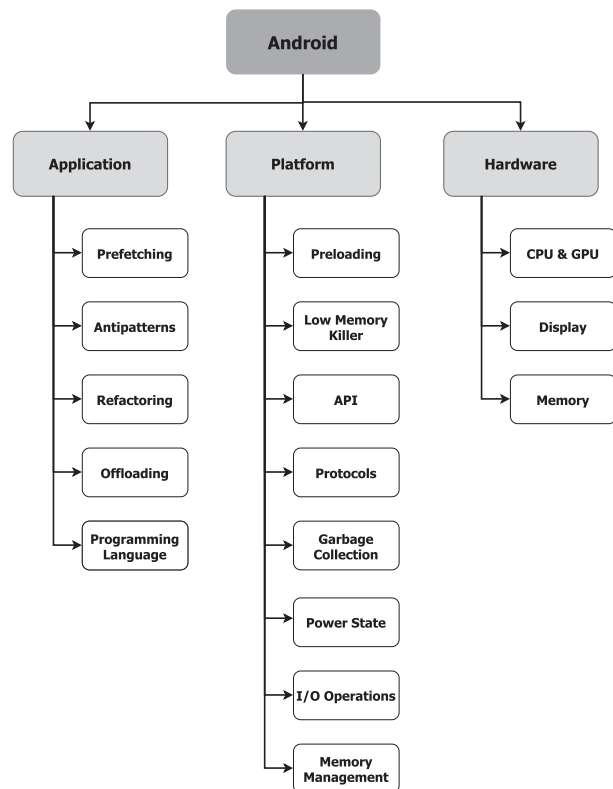


Fig. 1. Categorization of existing optimization approaches for non-functional performance characteristics of mobile applications.

The rest of this paper is structured as follows. Section 2 presents an overview of mobile devices and mobile-application ecosystems. The search methodology is described in Section 3. Sections 4, 5, 6, and 7 describe research on non-functional performance optimization. These refer to: responsiveness (Section 4), launch time (Section 5), memory consumption (Section 6) and energy consumption (Section 7). A discussion of results considering all non-functional performance optimization characteristics is given in Section 8. Section 9 presents related work and Section 10 outlines threats to validity. Section 11 concludes this survey.

## 2 BACKGROUND

This section presents an overview of the context of this survey. Initially, we present key terms and definitions regarding the architecture of mobile devices. Then, we focus on the characteristics of the Android platform that we take into account in this work. Finally, we explain how the function and performance of mobile applications can affect users.

### 2.1 Mobile Devices

Mobile devices are embedded systems that consist of hardware and software components. Fig. 2 illustrates a representative architecture of a mobile device.

The foundation of mobile devices is their hardware. The capacity of hardware components, such as physical memory, processors, and battery, is constrained. Additionally, mobile devices come with a growing set of embedded sensors, including accelerometers, digital compasses, GPS, microphones, and cameras, which enable the emergence of personal, group and community-scale sensing applications [23].

<sup>3</sup>. <https://gs.statcounter.com/os-market-share/mobile/worldwide>  
Authorized licensed use limited to: NANJING NORMAL UNIVERSITY. Downloaded on October 25, 2024 at 03:59:33 UTC from IEEE Xplore. Restrictions apply.

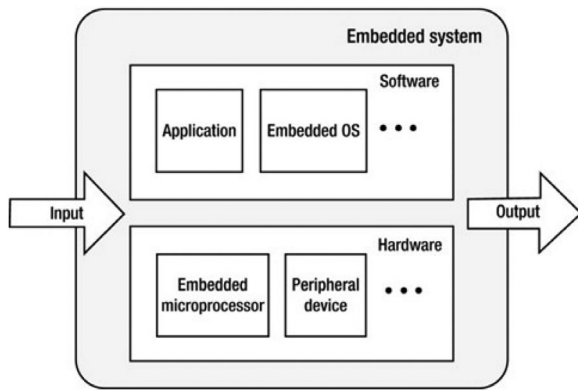


Fig. 2. Overview of mobile-device architecture [22].

However, the use of these sensors in applications requires a higher energy consumption [24].

The software of mobile devices comprises two basic layers: the mobile platform and the hosted mobile applications. The mobile platform (e.g., Android, ios) consists of an embedded operating system (os) that connects hardware with software components. It offers services such as memory management, networking, and power management. In mobile platforms, software libraries, which are used for the interaction with components, such as the database and media framework, are on the top of the os. Mobile applications (e.g., calculator, photos, contacts) are either provided by a mobile framework (e.g., the Android platform) or third-party applications provided by online stores for mobile applications (e.g., Google Play Store, Apple ios App Store, AppNokia, Samsung, BlackBerry World and Windows Phone Store).

## 2.2 Android

This survey focuses on the Android platform because it is currently the most used mobile platform<sup>4</sup> and open-source software, facilitating the analysis and evaluation of mobile systems. The following paragraphs present the main components of the Android platform.

Android is an embedded system based on the Linux os. The Linux kernel links hardware and software components of a mobile device. It manages services such memory, processes, power, and networking access, and it offers drivers for flash memory, Bluetooth, wifi, keyboard and audio. On top of the Linux kernel, lies the Android Runtime (ART), which is essential for running different applications. Each application runs as a separate process, having its own virtual machine instance.

The Android platform provides several methods and tools for improving the performance of mobile applications. For instance, memory is freed by the os if the available memory on a device is low. To achieve that, Android uses the Low Memory Killer (LMK) to remove the Least Recently Used (LRU) cached application from the memory. Cached data is stored in the virtual memory, as long as memory is available [25]. In order to optimize the cache memory, and

address problems such as duplicated pages in virtual memory, Kernel Same-page Merging (KSM) [26] and ZRAM [27] are applied by Android [28]. Even though these methods optimize the cache memory, they consume power while they are being executed.

Furthermore, Android offers a variety of tools in its SDK to analyze system information and support application development.<sup>5</sup> Tools can be used for logging (LOGCAT), retrieving application and system information (APKANALYZER, DUMPSYS, SysTRACE), as well as for simulations and debugging (ANDROID DEBUG BRIDGE, AVDMANAGER). Finally, Android provides developers with selected performance measures (Android Vitals)<sup>6</sup> that use real user data, in case users have agreed on providing such information. If that happens, several metrics related to startup time, battery usage, and crash stack traces are recorded. Such metrics can assist developers to monitor memory and energy consumption, to identify synchronization issues, and to avoid application crashes [29].

## 2.3 User Experience

User experience and satisfaction are important factors that can ensure the success of mobile applications [30]. Application rating, or user satisfaction with an application, has been shown to correlate with the number of downloads [31]. After installing and using an application, users are able to make judgements regarding their satisfaction. Reviews regarding user satisfaction of mobile applications appear in App Stores and new users consider them in order to decide whether they will download an application or not.

To achieve a high level of user satisfaction, developers focus on improving both the functional and non-functional characteristics of mobile applications [1], [2], [3], [13], [15], [16]. Apart from fatal issues with functionality, such as application crashes [3], non-functional performance characteristics also shape users' perception [1]. Non-functional performance characteristics are the first characteristics to, potentially adversely, affect users [15] and can lead to application uninstallations [2]. In the following, we describe functional and non-functional characteristics of mobile applications that can affect user experience.

*Functional* characteristics describe whether an application is doing what it is supposed to do (i.e., its behavior). Frequent complaints about functional aspects of applications include freezes or crashes [1], functional errors, such as not getting push notifications, and the removal of features [3].

*Non-functional characteristics* determine how an application carries out (performs) its behavior. Even though it is difficult to measure and judge non-functional characteristics [32], they represent a vital part of user satisfaction for mobile applications. Related work analyzes a range of different non-functional characteristics regarding applications' performance [1], [2], [3], [13], [15], [16].

Different schemes exist to classify functional and non-functional characteristics of applications [33], [34], [35]. Among these, the FURPS model [35], [36] clearly distinguishes *performance* characteristics from other functional and non-functional characteristics, as follows:

4. <https://gs.statcounter.com/os-market-share/mobile/worldwide/#monthly-201909-201909-bar>

5. <https://developer.android.com/studio/command-line>

6. <https://developer.android.com/topic/performance/vitals>



- **Functionality:** feature set, capabilities, generality, security;
- **Usability:** human factors, aesthetics, consistency, documentation;
- **Reliability:** frequency/severity of failure, recoverability, predictability, accuracy, mean time to failure;
- **Performance:** speed, efficiency, resource consumption, throughput, response time;
- **Supportability:** testability, extensibility, adaptability, maintainability, compatibility, configurability, serviceability, installability, localizability, portability.

Among these characteristics, we are interested in the *Performance* category, which in the context of mobile applications, can be further refined into the following four non-functional performance characteristics:

- 1) *Responsiveness* captures the time required to update the frames of the graphical user interface after user interaction.
- 2) *Launch time* describes the time required to start an application. This can happen as a cold start, when the application is launched without cached data (e.g., after device boot). Another launch type refers to the hot/warm start of application, which occurs when an application activity is kept in memory for a faster launch [37].
- 3) *Memory consumption* describes the amount of occupied memory. Memory can be shared between multiple applications or stored separately [38]. On constrained systems, such as smartphones, memory is a critical resource [25], [28].
- 4) *Energy consumption* is associated with the battery life. Energy is consumed by various components, including CPU, LCD, GPS, audio and wifi services [24], [39].

A detailed description of each of the four non-functional performance characteristics is given in Sections 4, 5, 6, and 7 (responsiveness, launch time, memory consumption, energy consumption, respectively).

### 3 SURVEY METHODOLOGY

The purpose of this survey is to gather and categorize research work published in the mobile computing and software engineering literature that refers to the optimization of non-functional performance of Android applications.

As this is an emergent topic and there is limited related work on Android performance optimization techniques to perform a systematic literature review (according to the guidelines of Kitchenham [40]), we conduct a *comprehensive literature review*. In the following, we present our search methodology in detail, starting with a preliminary and venue search, followed by a repository search and snowballing.

#### 3.1 Search Methodology

Our literature review on performance optimization includes publications that refer to optimization techniques on mobile applications and measurement of application performance.

##### 3.1.1 Preliminary Search

Prior to systematically searching online repositories, we conducted a preliminary search. The goal of the preliminary

search is to gain a deeper understanding of the field and assess whether there is a sufficient amount of publications that allows for subsequent analysis. Based on these results, we distinguish between four different non-functional performance characteristics: responsiveness, launch time, memory and energy consumption.

Other than Sadeghi *et al.* [41], who refined keywords during their search, we perform a preliminary search to guide our repository search. Additionally, we use the results of the preliminary search to define keywords (Table 2) and venues (listed in Section 3.1.2).

##### 3.1.2 Repository Search

Proceeding the preliminary search, we conduct a search of six established online repositories (IEEE, ACM, ScienceDirect, Scopus, arXiv, and Google Scholar). We have gathered publications from 2008 to February 2020, since the first version of Android was released in 2008.

To ensure that we provide an exhaustive literature search, we manually examine relevant venues from the field of software engineering and mobile computing, which we encountered during the preliminary search. We search venues with at least five publications in our preliminary search.

- *Conferences:* ICSE, ASE, MSR, MobiSys, MobileHCI, MobileSoft, UbiComp, CHI, ESEC/FSE.
- *Journals:* IEEE TSE.

##### 3.1.3 Selection

Table 2 lists keywords used to guide our repository search. Keywords are divided into five categories. First, the keywords that belong to the *Platform* category ensure that the selected publications deal with mobile platforms, particularly Android. Furthermore, keywords that belong to the *Responsiveness*, *Launch time*, *Memory*, and *Energy* categories filter publications referring to non-functional performance characteristics. We restrict search results to publications that contain at least one “platform” keyword for the platform and one keyword for “non-functional” performance in their title.

To ensure that the publications found during our search are relevant to the context of non-functional performance optimization of mobile applications, we consider the following inclusion criteria:

- The publication should refer to at least one of the non-functional performance characteristics investigated in this survey (e.g., responsiveness, memory, energy, and launch time), or to an approach that profiles at least one of the mentioned performance characteristics.
- The publication investigates the proposed methods on smartphones with an Android OS.

To assess whether the publications satisfy our inclusion criteria, we manually examined every publication using the process adopted by Martin *et al.* [42], as follows:

- 1) *Title:* First, all those publications whose title clearly does not match our inclusion criteria are excluded;
- 2) *Abstract:* Second, the abstract of every remaining publications is checked. Publications whose abstract

TABLE 1  
Results of the Repository Search

		Responsiveness	Launch Time	Memory	Energy		Responsiveness	Launch Time	Memory	Energy
Hits	GS	835	90	392	1129	IEEE	180	14	47	487
Title		31	10	16	76		21	5	9	52
Abstract		24	9	11	58		17	4	9	37
Body		11	1	4	31		9	0	7	8
Hits	Scopus	146	6	57	269	ACM	73	21	23	152
Title		20	3	6	51		17	10	3	47
Abstract		15	3	6	44		16	8	2	44
Body		7	0	2	15		8	4	2	20
Hits	SD	71	13	97	62	arXiv	17	1	1	18
Title		2	2	0	21		3	0	0	10
Abstract		2	0	0	15		1	0	0	9
Body		0	0	0	1		2	0	0	2

The number of papers retained at each stage of the search (e.g., Hits, Title, Abstract, Body) is given for each online repository (e.g., Google Scholar, IEEE, Scopus, ACM, Science Direct, arXiv) and non-functional performance characteristic (Responsiveness, Launch Time, Memory, Energy). Google Scholar is abbreviated with GS; Science Direct is abbreviated with SD. For example, searching for “responsiveness” in GS retrieves 835 publications, among those 31 have a relevant title, 24 of those have an abstract satisfying our inclusion criteria, and a total of 11 (of those 24) publications are included in our survey after reading them entirely.

TABLE 2  
Keywords Used for the Repository Search

Category	Keywords
Platform	android, smartphone, app, apps
Responsiveness	responsiveness, performance
Launch time	launch, start
Memory	memory
Energy	energy, battery, power

does not meet our inclusion criteria are excluded at this step;

- 3) *Body*: Publications that passed the previous two steps are then read in full, and excluded if their content does neither satisfy the inclusion criteria nor contribute to this survey.

Based on the above three-stage process and inclusion criteria, we iteratively reduce the amount of publications obtained from online repositories, until we end up with the set of publications investigated in the following sections (Sections 4, 5, 6, and 7). This process is performed by two authors independently, the results are compared at each stage, and disagreements discussed until an agreement is reached.

### 3.1.4 Snowballing

After a collection of publications is obtained from the repository search, we proceed to inspect the related work of the publications selected in the previous search to gather cited publications using snowballing [43]. We apply one level of backwards snowballing.

## 3.2 Selected Publications

Table 1 shows the results of the repository search. The amount of publications found during each step of the search is listed.<sup>7</sup>

7. A collection of publications after checking *Abstract* is available in our on-line appendix [21].

In the following, we give the number of unique publications after each stage of the search procedure in addition to the number of newly added publications:

- 1) Preliminary search: 96
- 2) Repository search: 174 (+80)
- 3) Venue search: 180 (+4)
- 4) Snowballing: 252 (+72)
- 5) Author feedback: 297 (+45)

In addition to the discussed stages of the search procedure (1-4), we added 45 publications based on the feedback from the authors cited. Among all 294 publications, 156 unique publications optimize at least one non-functional performance characteristic. These 156 publications were published in 97 different venues. We further classify top publication venues (A, A\* according the CORE ranking Portal),<sup>8</sup> regarding their category based on the ACM’s Computing Classification System (ccs).<sup>9</sup> Among these, a majority of publications is obtained from Software Engineering (32 percent), and Computer Systems Organisation (29.33 percent) venues. The remaining publications are retrieved from Mobile Computing (20 percent), Networks (17.33 percent), and Security and Privacy (1.33 percent) venues. A full list of conferences and journals is available online [21].

The publication distribution over the entire search period is illustrated in Fig. 3. Note, a publication can contribute to more than one subtotal if it explicitly optimizes more than one non-functional performance characteristic. During our search, we found ten publications that optimize more than one non-functional performance aspect [15], [37], [44], [45], [46], [47], [48], [49], [50], [51]. Among these, there is one publication that optimizes three characteristics (responsiveness, energy consumption, memory consumption) [44], while the others optimize two. Section 8.3 provides further details on the relationships between performance characteristics.

8. <http://www.core.edu.au/conference-portal>. Additionally, we include MobiSys, classified as “B”, due to its popularity on mobile systems.

9. <https://dl.acm.org/ccs>




Fig. 3. Number of publications on non-functional performance optimization per year.

Based on our search results, we devise the categorization of approaches shown in Fig. 1. These categories consist of approaches (e.g., offloading, code optimization) and elements of the Android platform (e.g., Low Memory Killer, API). In the remainder of the survey, we discuss our search results in detail.

## 4 RESPONSIVENESS

Responsiveness refers to the ability of mobile applications to respond to user interactions fast and smoothly. An application is highly responsive when the time it takes to respond to user requests is minimal. A highly responsive application offers high user satisfaction as users prefer not to wait when interacting with an application. On the other hand, an application with poor responsiveness can have a negative impact on user perception, and on its success [52].

Specifically, Tolia et al. [53] argued that response times lower than 150 ms do not negatively affect user satisfaction. In fact, delays that last almost one second do not significantly affect users, but start making them aware of these delays, whereas delays that last more than one second indeed make users “unhappy” [53]. Willocx et al. [54] stated that response times under 100 ms appear as instantaneous to users. Furthermore, users would accept response times up to a few seconds if delays were to occur rarely [54].

To detect and fix hot spots in mobile applications that may cause response-time delays, developers use several techniques, including profiling and optimization approaches. Section 4.1 presents methods and approaches for profiling, while Section 4.2 presents approaches to optimize responsiveness. Section 4.3 summarizes our findings.

### 4.1 Profiling

There are several profiling approaches that developers use to measure the responsiveness of applications and locate hot spots for improvement. Popular profiling techniques include the measurement of page loading time [55], the measurement of the overall frame time and the calculation of the number of delayed frames [50], as well as the estimation of Central Processing Unit (CPU) time [56]. Furthermore, responsiveness can be measured at different levels of the Android platform, considering the UI [57] and hardware components [58].

Several tools have been developed to measure the responsiveness of mobile applications. Specifically, Ravindranath et al. [59] introduced APPINSIGHT to detect critical paths in applications, which represent bottlenecks for user transactions. Hong et al. [60] proposed PERFPROBE, a profiling

TABLE 3  
Studies on Responsiveness Optimization

Category	Authors [Ref]	Year	Venue
Offloading	Kemp et al. [66]	2010	MobiCASE
	Chun et al. [67]	2011	EuroSys
	Ra et al. [68]	2011	MobiSys
	Kosta et al. [69]	2012	INFOCOM
	Gordon et al. [70]	2012	OSDI
	Gordon et al. [71]	2015	MobiSys
	Das et al. [45]	2016	IACC
	Montella et al. [72]	2017	CCPE
	Chen and Hao [73]	2018	J-Sac
Antipatterns	Jin et al. [74]	2012	SIGPLAN
	Yang et al. [52]	2013	MOBS
	Nistor et al. [75]	2013	ICSE
	Liu et al. [13]	2014	ICSE
	Ongkosit and Takada [76]	2014	DeMobile
	Hecht et al. [77]	2015	ASE
	Habchi et al. [78]	2018	ASE
	Hecht et al. [50]	2016	MobileSoft
	Li et al. [79]	2019	SANER
Refactoring	Lin et al. [80]	2014	FSE
	Okur et al. [81]	2014	ICSE
	Lin et al. [82]	2015	ASE
	Lyu et al. [46]	2018	ISSTA
	Feng et al. [83]	2019	ICSTW
Prefetching	Higgins et al. [84]	2012	MobiSys
	Zhao et al. [85]	2018	ICSE
	Choi et al. [86]	2018	CoNEXT
	Malavolta et al. [87]	2019	ICSE-NIER
Programming languages	Batyuk et al. [88]	2009	MobileWare
	Lee and Jeon [89]	2010	ICCAS
	Lee and Lee [90]	2011	iCast
	Lin et al. [91]	2011	IBICA
	Saborido et al. [44]	2018	EMSE
CPU & GPU	Wang et al. [92]	2013	CGO
	Cheng et al. [93]	2013	IWSSIP
	Thongkaew et al. [94]	2015	JIP
I/O operations	Nguyen et al. [49]	2015	MobiSys
	Mao et al. [95]	2018	ITCSDI
Hardware components	Kim and Shin [96]	2015	ICUIMC

approach to diagnose hardware and software causes for slowdowns with runtime information. Kim et al. [61] conducted performance testing, using unit tests, at early development stages of the applications to identify response-time delays. Kang et al. [62], [63] presented a technique that analyzes application performance focusing on particular asynchronous executions. Wang and Rountev [64] introduced a novel approach that profiles responsiveness by tracking the usage of mobile resources such as bitmap or SQLite databases. Kwon et al. [65] proposed MANTIS, a framework that predicts the execution time of an application while using particular inputs.

### 4.2 Optimization Approaches

For a high responsiveness of mobile applications, developers apply several categories of optimization approaches. In the following, we describe techniques found in literature. Table 3 lists our findings.



Offloading refers to the transfer of heavy computational tasks to external computing units with less performance-related restrictions. This technique is popular in the development of mobile applications as mobile devices are embedded systems with restricted memory and CPU. However, offloading comes with an overhead while transferring the results of processes from an embedded system to external computing units and vice versa [71]. The first offloading implementation for Android applications was introduced by Kemp *et al.* [66] and refers to the CUCKOO framework. CUCKOO helps developers to easily implement offloading tasks in their applications. The offloading decision is made at runtime based on heuristics, context, and historic information. Other frameworks that support offloading include CLONECLOUD [67], THINKAIR [69], COMET [70], and TANGO [71]. Offloading has frequently been used for responsiveness improvements [45], [68], [72], [73].

Antipatterns are bad programming patterns, such as performance bugs, which deteriorate software quality and reduce application responsiveness that can negatively affect user experience [13], [50]. For this reason, several tools have been developed for the identification and removal of antipatterns. In particular, Liu *et al.* [13] introduced PERFChecker to automatically detect performance bugs in mobile applications. PERFChecker is built on top of the SOOT [97] Java optimization framework and analyzes applications at a bytecode level. PERFChecker applies static code analysis to detect antipatterns. Other tools focus on detecting particular types of bugs related to application responsiveness [50], [75]. For instance, Nistor *et al.* [75] searched for repetitive computations in code loops, following the intuition that repetitive behavior is likely to be optimizable. To detect repetitive behavior, they created TODDLER, an automated oracle to analyze memory access patterns. Hecht *et al.* [50] used a static analysis tool called PAPRIKA [77] to detect three types of code smells (Internal Getter/Setter, Member Ignoring Method, and HashMap Usage). They investigated the removal of code smells in an empirical study, obtaining responsiveness improvements of up to 12.4 percent.

Furthermore, Inefficient Image Displaying (IID) can cause performance degradation (e.g., repeated and redundant image decoding). Li *et al.* [79] developed the static analysis tool TAPIR to detect IID issues, which can be strongly correlated with antipatterns.

Finally, responsiveness-related bugs can be detected using predefined rule sets (e.g., efficiency rules) [74], [76] and test amplification (insertion of artificial delays in application source code) [52].

Refactoring can be performed to utilize efficient programming practices. Lin *et al.* [80], [82] provided an analysis showing that even though applications include concurrent code, they often contain bugs or end up with executing the source code sequentially. For concurrent code execution, and higher responsiveness, the authors located and refactored long-running operations, using the two tools ASYNCHRONIZER and ASYNCDROID. Okur *et al.* [81] developed two tools to refactor asynchronous code in Windows Phone applications. Asynchronous code is converted by ASYNCIFIER and common misuses in asynchronous code are corrected by CORRECTOR. Their empirical study showed that developers accept the proposed changes to asynchronous code. Lyu *et al.* [46] applied static analysis to change inefficient

database operations that are placed in within a loop. Database operations called in loops can cause Repetitive Auto-commit Transaction (RAT), which creates a new transaction in each iteration of the loop. Furthermore, Feng *et al.* [83] mined optimization patterns from GITHUB projects, considering performance-aware APIs, which can be manually injected into the source code of mobile applications and improve their performance.

Prefetching refers to a technique that caches data in advance, so that it can timely provide required data when it is needed. Higgins *et al.* [84] provided a library called Informed Mobile Prefetching (IMP) that assigns the task of determining when to prefetch data to the mobile system, rather than leaving the choice to developers. Developers solely specify which items could benefit from prefetching, while IMP determines whether and how prefetching is handled, based on responsiveness, battery lifetime, and mobile data usage. Zhao *et al.* [85] proposed a technique named PALOMA, that prefetches HTTP requests to reduce responsiveness latency. PALOMA uses string analysis to detect prefetchable content in the application source code. While users navigate in an application, PALOMA uses short pauses (“user think time”) for prefetching. Choi *et al.* [86] identified resource dependencies with static analysis to automatically generate acceleration proxies for dynamic prefetching. Application binary files are analyzed to detect HTTP(s) messages, which are later used for prefetching. As static analysis lacks certain information, missing information of HTTP(s) requests is added at runtime. Lastly, Malavolta *et al.* [87] proposed a technique called NAPPA that prefetches network requests based on user navigation patterns.

Programming languages can impact the processing speed of applications and therefore responsiveness. Android provides the Native Development Kit (NDK) that allows developers to write native C/C++ code. Native instructions are directly executed by the CPU and, therefore, they provide a better performance over non-native ones [92], [98]. Several empirical studies compared the performance of programming languages, and found that native C code reduces the running time of the same algorithms written in Dalvik Java code [88], [89], [90], [91]. Furthermore, efficient implementation choices, such as which map variant to use (e.g., HashMap, ArrayMap, and SparseArray) can improve responsiveness [44].

CPU and Graphics Processing Unit (GPU) adaptations can accelerate the execution of time-consuming programming tasks and increase application responsiveness. Wang *et al.* [92] proposed ACCELDROID to accelerate the execution of bytecode on the HW/SW co-designed processor of Android. Therefore, instead of translating bytecode twice, this is only translated once. Cheng *et al.* [93] provided guidelines to map applications to the Android platform (e.g., whether to use CPU or GPU and how many cores are used). This mapping is platform as well as task-dependent. An optimal performance choice can avoid performance degradation. Additionally, Thongkaew *et al.* [94] developed architectural hardware extensions that can fetch and decode Dalvik bytecode directly.

I/O operations have an impact on responsiveness and can enable optimization [99]. For instance, Nguyen *et al.* [49] proposed an approach that adapts the prioritization of read

and write operations for avoiding slowdowns. Mao *et al.* [95] introduced a trace collection tool to identify redundant I/O requests in mobile applications and eliminate them to reduce response times. As redundancy is minimally shared among applications, they performed an application-aware optimization.

*Hardware components*, such as the use of embedded Multi-media Cards (eMMC), can be investigated for responsiveness improvements. Kim and Shin [96] studied whether additional features of eMMCs are utilized by Android smartphones, and reduced the I/O latency.

### 4.3 Summary

Responsiveness is a non-functional performance characteristic concerned with the time an applications needs to respond to user requests. In practice, this is either measured in time (ms) [55], [56] or in frames [50]. Several tools have been proposed to measure responsiveness and detect responsiveness issues [59], [60], [62], [63].

Since responsiveness measures the duration required to complete computations, a naive approach to improve response times is to move the computations from the smartphone to devices with less restrictions. This approach is called “offloading” and requires additional infrastructure (e.g., external servers for computation). If such an infrastructure is not available, other approaches can be followed, which in majority are applicable to a source-code level (23/38).

Furthermore, changes that have been applied to mobile applications’ source code include the removal of bad programming patterns (antipatterns), and the usage of good programming practices (e.g., concurrency with the help of refactoring [80], [81], [82], and prefetching of content [84], [85], [86], [87]). Other than changing source code after responsiveness issues have been detected, a carefully considered choice of the right programming language can lead to improvements (e.g., native C/C++ is faster when executed on CPUs [88], [89], [90], [91]). Lastly, changes to the hardware have achieved responsiveness improvements. This can directly happen at CPU or GPU level [92], [93], which can carry out computations, or on other hardware components (e.g., memory [96]).

## 5 LAUNCH TIME

During the launch of a mobile application, operations and data are loaded to make the application available to the user. Therefore, launch time is the first performance characteristic of a mobile application that users have the opportunity to notice. Launch time directly influences user experience and satisfaction. Nagata *et al.* [100], Song *et al.* [37] and Kim *et al.* [15] defined the launch time of an application as the required time until user input is accepted. Yan *et al.* [101] described the Total Launch Time (TLT) of an application as the needed time until the entire content, including asynchronously loaded content, can be displayed to the user.

Furthermore, Song *et al.* [37] found that cold start time (when an application is started from scratch) has a significant impact on the application launching experience of users. Developers can address this issue by analyzing their applications’ source code to identify and fix bottlenecks that

TABLE 4  
Studies on Launch Time Optimization

Category	Authors [Ref]	Year	Venue
Preloading	Yan <i>et al.</i> [101]	2012	MobiSys
	Nagata <i>et al.</i> [100]	2013	CANDAR
	Parate <i>et al.</i> [106]	2013	UbiComp
	Tang <i>et al.</i> [107]	2013	SIGAPP
	Chung <i>et al.</i> [47]	2013	TECS
	Song <i>et al.</i> [37]	2014	TECS
	Lee <i>et al.</i> [48]	2017	J-SAC
	Baumann and Santini [108]	2017	IMWUT
	Martins <i>et al.</i> [16]	2018	ICMLT
Low Memory Killer	Chung <i>et al.</i> [47]	2013	TECS
	Prodduturi and Phatak [109]	2013	IIT
	Song <i>et al.</i> [37]	2014	TECS
	Baik and Huh [110]	2014	ICSE
	Kim <i>et al.</i> [15]	2015	IEEE Micro
	Vimal and Trivedi [111]	2015	RAICS
	Singh <i>et al.</i> [112]	2016	IOTA
	Kim <i>et al.</i> [103]	2016	TECS
	Lee <i>et al.</i> [48]	2017	J-SAC
	Li <i>et al.</i> [113]	2017	IWCMC
Memory	Joo <i>et al.</i> [104]	2011	FAST
	Kim <i>et al.</i> [15]	2015	IEEE Micro
I/O operations	Nguyen <i>et al.</i> [49]	2015	MobiSys

possibly increase launch time. The following sections discuss profiling methods (Section 5.1) and approaches to optimize launch time and cold start issues (Section 5.2). A summary is given in Section 5.3.

### 5.1 Profiling

Profiling approaches have been proposed for locating issues in mobile applications that may increase the application launch time. Using monitoring functions in the source code of Android applications, in the Android platform, and in third-party libraries used by Android applications, developers can pinpoint performance issues causing launch-time delays [100], [102]. Also, developers can profile the usage of system resources to pinpoint the application launch completion [103]. Additionally, Nguyen *et al.* [49] studied launch delays of an application as the time taken in kernel mode and the time spent waiting for disk network operations.

To understand how launch time can affect user behavior, Song *et al.* [37] investigated logs from application usages. Other approaches for assessing the launch behavior of applications include monitoring of: handling of I/O requests [104], system memory usage [15], restart ratio of applications (the number of cold starts over all application launches) [37] and user satisfaction with regards to launch-time delays [105].

### 5.2 Optimization Approaches

In order to reduce launch time, developers apply optimization techniques. Table 4 summarizes publications found in literature, which are described in the following.



Preloading of application data prevents cold starts and thereby the overall user waiting time. Frequently, the usage of applications follows patterns [37], and enables the prediction and preloading of the next application to be used. Context information, such as the time of day or the location can be taken into account to improve predictions [16], [48], [114] and preload applications.

The FALCON approach by Yan *et al.* [101], which refers to an OS extension, preloads applications and application-specific content based on the context (e.g., location) and usage patterns. For this purpose, spatial and temporal features are designed based on an extensive analysis. Usage patterns include the use of weather applications in the morning, or playing games at home.

Additionally, application predictions determine the applications to be launched next [37], [47], [107], [108] and when they are going to be used [106]. In particular, Nagata *et al.* [100] analyzed the relationship of application launch time regarding the number of preloaded classes. For this, they manually selected a number of preloaded classes, and showed for one application that the launch time is reduced when the number of preloaded classes is high. The prediction of the next application to be used has been also investigated for restructuring user interfaces [115], [116], [117].

The Low Memory Killer (LMK), which removes the LRU application from memory, may not lead to optimal results [15], [112], because users do not always rely on recently used applications. To address this issue, several studies introduce techniques that determine which data should be removed from memory. Specifically, Song *et al.* [37] and Li *et al.* [113] devised models to detect patterns in application usages based on application cold start times. Decisions for the LMK are based on usage patterns to prioritize data of applications that are likely to be launched.

Instead of removing the LRU application from memory, choices can be made based on application cold start times [15], the required storage size [103], and the importance of an application to a user [109], [111], [112]. Furthermore, Baik and Huh [110] analyzed usage patterns and determined a threshold on how many processes to keep in memory before freeing them. If this limit is fixed to a high value, more applications can be kept in memory, leading to fewer restarts.

Memory can be adapted to suit application launches better. Accordingly, Joo *et al.* [104] proposed the use of SSDs instead of HDDs to speedup application launch time. This approach was not designed for mobile devices; however, mobile devices use NAND flash memory as secondary storage carrying almost identical performance characteristics as SSDs. One could therefore apply this approach to mobile devices as well. Furthermore, Kim *et al.* [15] proposed the use of Non-Volatile Memory (NVM) to store frequently used applications and shared libraries among applications. Shared data is stored on Phase-Change Memory (PCM). Therefore, less data needs to be loaded when launching applications.

I/O operations impact the application launch time, as their speed can be seen as a performance bottleneck during application launch [15]. In fact, Nguyen *et al.* [49] analyzed the impact of read and write operations on launch time, as mobile devices wait for I/O operations to complete. As

application launches are dominated by read operations (five times as many read operations as write operations [49]) this can have a high impact on overall application launch time. A prioritization of read and write operations avoids slowdowns and reduces launch time.

### 5.3 Summary

Launch time describes the time required until a user input is received [15], [37], [100] or the entire application content is displayed [101], after an application has been started by the user. The application launch completion has been profiled according to the system's resource-management usage [103].

Due to the high negative impact that cold starts can have on user satisfaction [37], the majority of launch time optimization methods (15/18) prevent cold starts, and reduce application launch time. On one hand, preloading of application data can be applied to spend loading times before the application launch, and reduce the actual launch itself. For this purpose, predictions are used to determine applications that are likely to be used next, based on usage patterns [16], [37], [48], [101], [114]. On the other hand, changes to memory management (LMK) have achieved similar results. In contrast to preloading, which loads desired application data, changes to the LMK to keep important data in memory for a longer time. Both of these approaches require access to the Android OS to implement the required adaptations.

An increased speed of memory operations (e.g., usage of SSDs over HDDs [104]), shared libraries among applications [15], and I/O prioritization of read over write operations [49] have been also applied to reduce launch time.

## 6 MEMORY

Memory is a critical resource for embedded systems [25], [28], such as mobile devices. Specifically, in Android devices, data can be loaded either by the Android platform (to be shared across multiple running applications) or by each application, separately. Application data is stored in separate heaps, per application [38]. The main memory is typically shared between the CPU and GPU [51]. Therefore, a considerable amount of the main memory is occupied by graphic processing operations [25]. Kim *et al.* [15] classified applications in two categories, based on their memory consumption, *stable* and *unstable*. The memory consumption of *stable* applications increases within the first ten seconds of the applications' launch, and it stabilizes afterwards. The memory consumption of *unstable* applications increases steadily, and it does not stabilize.

The following sections present approaches that are used to measure the consumed memory of mobile applications (Section 6.1), optimize application memory consumption (Section 6.2). At the end, we provide a summary (Section 6.3).

### 6.1 Profiling

Different tools and approaches have been proposed to measure the memory usage of applications. For instance, memory consumed by mobile applications can be measured by: kernel memory footprints [118], garbage collection calls [50], physical memory dumps, and logging information [28]. Vimal and Trivedi [111] used the Dalvik Debug Monitor

TABLE 5  
Studies on Memory Optimization

Category	Authors [Ref]	Year	Venue
Antipatterns	Park and Choi [121]	2012	IJCA
	Shahriar <i>et al.</i> [122]	2014	HASE
	Hecht <i>et al.</i> [50]	2016	MobileSoft
	Santhanakrishnan <i>et al.</i> [123]	2016	i-Society
	Tasneem <i>et al.</i> [51]	2019	IJCA
	Amalfitano <i>et al.</i> [124]	2020	IEEE Access
Garbage collection	He <i>et al.</i> [38]	2011	IFIP
	Gerlitz <i>et al.</i> [125]	2013	JTRES
	Lim <i>et al.</i> [126]	2013	ICCE
	Mori <i>et al.</i> [127]	2017	GCCE
	Tasneem <i>et al.</i> [51]	2019	IJCA
Deduplication	Kim <i>et al.</i> [128]	2014	ICSE
	Lee <i>et al.</i> [28]	2015	APSys
Memory management	Kim <i>et al.</i> [129]	2013	IEEE TCE
	Jeong <i>et al.</i> [130]	2013	USENIX ATC
	Zhong <i>et al.</i> [131]	2014	EMSOFT
	Kim <i>et al.</i> [132]	2015	ISMM
	Nguyen <i>et al.</i> [133]	2016	WiMob
	Kim and Bahn. [134]	2017	PMC
	Kim <i>et al.</i> [135]	2017	TECS
	Kim and Bahn. [136]	2019	IEEE Access
GPU	Kwon <i>et al.</i> [25]	2015	EMSOFT
Programming languages	Escobar De La Torre and Cheon [137]	2017	UTEP
	Saborido <i>et al.</i> [44]	2018	EMSE

Server (DDMS) to analyze memory footprints of Android components and measure memory consumption.<sup>10</sup> Tools such as ANDROSCOPE by Cho *et al.* [58] have been also used to analyze the performance (including memory) of all the layers of the Android platform. Furthermore, ANDROBENCH [119] and ANDROSTEP [120] are benchmark tools that assess the storage performance of Android devices by analyzing logs from read and write I/O operations.

## 6.2 Optimization Approaches

To reduce the memory usage of mobile applications developers use different categories of optimization approaches. The following paragraphs summarize the relevant approaches found in literature. Table 5 lists the representative studies.

*Antipattern* coding practices can be used to identify code that is likely to lead to memory leaks. Memory leaks occur when applications constantly request memory while running [121], or when unused objects are being kept in memory longer than required [122].

Hecht *et al.* [50] showed in an empirical study that memory consumption can be reduced by correcting code smells. In particular, memory can be improved in terms of memory usage and number of garbage collection calls. Shahriar *et al.* [122] developed memory leak patterns for Android applications and used fuzz testing to emulate and detect memory leaks. A total of three fuzzing types (application, resource, and API) are used in their experiments, which discovered

crashes due to memory leaks in real-world applications. Furthermore, memory leaks can be identified by analyzing memory dumps [51], [138], the activity lifecycle [124], source code patterns [123] or memory execution information by applying process control block hooking [121].

*Garbage collection* is used in Android to manage memory and identify unused objects that can be removed [51]. Since version 2.2., Android uses a stop-the-world (stw) garbage collector [38], [125]. This stops other operations to free the memory and resume them afterwards, resulting in pauses that can negatively effect user experience [38].

Different garbage collector designs have been evaluated for improvements: reference counting garbage collection [51], [125], concurrent garbage collection [38] and generational garbage collection [38], [127]. Lim *et al.* [126] proposed a memory partitioning scheme, which partitions available memory into two nodes (for critical and uncritical applications). If one node runs out of memory, only the memory of this node is freed.

*Deduplication* is a technique to remove redundant pages from memory. While duplicated memory reduces available memory for other applications, Android is prone to have page-level duplication in memory [28]. Lee *et al.* [28] developed a system (MEMSCOPE) to analyze memory duplication in Android OS. MEMSCOPE identifies memory segments that contain duplicated memory pages. One of the disadvantages of deduplication is the additional computation needed to detect and merge redundant pages. Therefore, Kim *et al.* [128] proposed a computationally efficient deduplication scheme, considering background applications that do not update memory contents and need to be scanned only once.

*Memory management* changes can be applied to achieve further improvements in memory usage by mobile applications. For example, *swapping* is a technique that reclaims memory by writing inactive memory pages to secondary storage (e.g., eMMC). Kim *et al.* [135] proposed a swapping scheme (Application-Aware Swapping) that considers OS processes in the swapping decision. For example, swapping an application to secondary memory is not useful if the LMK is about to remove it from memory, freeing the used memory pages. Other approaches utilized NVM for swapping [131], [136].

Journaling in Android applies a write-twice behavior, to ensure reliability, which reduces system performance by additional write operations. Kim *et al.* [129] proposed an architecture to reduce storage accesses for journaling. They use non-volatile memory for this purpose. Among others, Jeong *et al.* [130] eliminated the journaling of unnecessary metadata. Nguyen *et al.* [133] proposed *iram*, a system that cleans low-priority processes to maintain a high level of free memory. Kim and Bahn [134] evicted write-only-once data from the buffer cache to improve the utilization of cache space. Kim *et al.* [132] proposed an approach to group memory pages with the same lifetime to alleviate fragmentation of I/O buffers.

GPU buffers have been analyzed by Kwon *et al.* [25], who introduced a compressing scheme. Once an application goes to the background, its GPU buffers are treated as *inactive* and compressed. If the application is launched in the foreground, GPU buffers are decompressed.

10. DDMS is deprecated and was removed from Android Studio 3.2. Android offers other tools to carry out the functions of DDMS <https://developer.android.com/studio/profile/monitor>.

Programming languages influence the choice of language constructs that further impact storage requirements. Escobar De La Torre and Cheon [137] analyzed the impact of the Java language constructs on the allocated memory. For instance, for-each loops require more memory than equivalent code snippets using regular loops. Removing those constructs (iterators, for-each loops, lambda expressions and the Stream API) reduces memory requirements [137]. Saborido *et al.* [44] showed that map implementations consume different amounts of memory. Specifically, `ArrayMap` uses less memory than `HashMap`.

### 6.3 Summary

Memory describes the occupation of device memory by applications, and is critical for resource-constraint systems [25], [28]. For Android applications, data is either shared between multiple applications, or loaded separately by each application.

Memory has been measured according to kernel memory footprints [118], garbage collection calls [50], physical memory dumps, and logging information [28]. Memory can be analyzed by tools provided by the Android OS [111], and external tools provided by researchers [58], [119], [120].

Memory consumption has been reduced by removing code smells from application source code [50]. In particular, memory leaks (e.g., constantly requesting memory [121], or keeping unused objects in memory [122]) have a negative impact on memory consumption.

Memory consumption has been further improved by changes in the Android OS. For example, garbage collection, which is used to free memory in Android, can use different strategies for freeing memory [109]. Another approach is the removal of redundant data from memory (deduplication) [28], [128]. Improvements can furthermore be achieved by changes in swapping [131], [135], [136] and journaling strategies [129], [130].

## 7 ENERGY

Embedded systems include several components that consume battery. CPU, LCD, GPS, audio and WiFi services are power-intensive components [24], [39]. Due to the limitation in battery size and stored energy [139], [140], [141], reducing energy consumption is gaining more and more relevancy [142]. In general, optimizing energy consumption depends on individual usage [7], [143].

The following sections outline methods and tools to profile energy consumption (Section 7.1) and reduce the energy consumption of mobile applications (Section 7.2). A summary is given in Section 7.3.

### 7.1 Profiling

Several measurements and prediction approaches have been used to profile energy consumption on mobile devices. Hoque *et al.* [144] discussed two ways to measure energy consumption: with external instruments and self-metering. This section gives an overview of respective profiling techniques.

A common approach to determine energy consumption is to investigate hardware components. Zhang *et al.* [24] measured power consumption using battery voltage sensors

and knowledge of battery discharge behavior. Additionally, fuel gauge chips [145] and Battery Monitoring Unit [146] can be used to measure energy consumption. Other approaches use physical power meters to measure energy consumption [98], [147], [148], [149], [150]. Morales *et al.* [142] used a digital oscilloscope for high frequency energy measurements. Ferrari *et al.* [151] designed a Portable Open Source Energy Monitor (POEM) to measure energy consumption of applications at a control flow level. Bokhari *et al.* [152] built energy models based on CPU utilization and lines of code, as external meters can be expensive and not easy for developers to set these meters up.

Other than measuring energy consumption with physical devices, several studies make energy consumption estimates. Energy consumption estimates can be performed based on hardware utilization and system-calls [153], Android kernel monitoring [154], pixel information [155], [156], user behavior [7], data transmission-flow characteristics [157], code level [158], [159] and source-code line level [160].

Jabbarvand *et al.* [161] proposed COBWEB, a search-based technique to generate test suites for energy testing. These tests are able to execute energy-greedy parts of the code. The computational cost of such a testing technique can be reduced by test-suite minimization [162]. Mittal *et al.* [163] presented an emulation tool WATTSOn to estimate energy consumption during application development. CPU time has been used as a proxy for energy consumption. However, it is not as accurate as other techniques, because voltage is scaled dynamically and multiple hardware components are used [160]. One should consider that errors during the measurement and estimation of energy consumption, as noise, can be introduced by various hardware components, such as a rising temperature of the battery [164]. This impacts the number of samples required for ensuring statistical significance when comparing the energy consumption of applications [165]. Validation approaches should consider the level of noise to compare solutions fairly [166].

### 7.2 Optimization Approaches

Several techniques have been applied to reduce the energy consumption of mobile applications, which are discussed in the following and summarized in Table 6.

*Offloading*, e.g., transferring computationally expensive tasks to external devices, can be used to reduce energy consumption [148]. Cuervo *et al.* [148] developed MAUI, a system that supports automatic and developer-specified code offload. MAUI determines which method to execute remotely based on the current state of the device at runtime. Offloading decisions can be motivated by device status [174], execution times [170], network conditions [45], [167], [168], [169] or developer decisions [171]. Bolla *et al.* [173] proposed the concept of Application State Proxy (ASP) to offload entire applications. ASP transfers internet-based applications to other network devices, when they are kept in the background. As long as no new events occur (e.g., messages), applications are kept in the proxy, which reduces the resource load on the smartphone. Corral *et al.* [172] applied offloading to matrix multiplication and image processing tasks to reduce energy consumption.



TABLE 6  
Studies on Energy Optimization

Category	Authors [Ref]	Year	Venue
Offloading	Cuervo <i>et al.</i> [148]	2010	MobiSys
	Saarinen <i>et al.</i> [167]	2012	SIGCOMM
	Ding <i>et al.</i> [168]	2013	SECON
	Saarinen <i>et al.</i> [169]	2013	Mobicom
	Khairy <i>et al.</i> [170]	2013	IWCMC
	Kwon and Tilewich [171]	2013	ICSME
	Corral <i>et al.</i> [172]	2014	MobiWis
	Bolla <i>et al.</i> [173]	2014	NGMAST
	Qian and Andresen [174]	2015	IJNDC
	Das <i>et al.</i> [45]	2016	IACC
Prefetching	Balasubramanian <i>et al.</i> [175]	2009	SIGCOMM
	Chen <i>et al.</i> [176]	2013	SOSP
	Mohan <i>et al.</i> [177]	2013	EuroSys
	Yang and Cao [178]	2017	IEEE TCM
	Dutta and Vandermeer [179]	2017	TWEB
Antipatterns	Pathak <i>et al.</i> [180]	2011	HotNets
	Zhang <i>et al.</i> [140]	2012	CODES
	Pathak <i>et al.</i> [181]	2012	MobiSys
	Banerjee <i>et al.</i> [147]	2014	FSE
	Liu <i>et al.</i> [182]	2014	TSE
	Jabbarvand and Malek [183]	2017	FSE
Refactoring	Pathak <i>et al.</i> [184]	2012	EuroSys
	Anwer <i>et al.</i> [185]	2014	MobileSoft
	Alam <i>et al.</i> [186]	2014	DATE
	Li <i>et al.</i> [187]	2014	ICSE
	Linares <i>et al.</i> [155]	2015	FSE
	Bruce <i>et al.</i> [188]	2015	GECCO
	Cito <i>et al.</i> [145]	2016	ASE
	Banerjee and Roychoudhury [189]	2016	MobileSoft
	Cruz <i>et al.</i> [190]	2017	MobileSoft
	Banerjee <i>et al.</i> [191]	2017	TSE
	Morales <i>et al.</i> [142]	2017	TSE
	Cruz and Abreu [192]	2017	FSE
	Bokhari <i>et al.</i> [164]	2017	GECCO
	Cruz and Abreu [193]	2018	CibSE
	Lyu <i>et al.</i> [46]	2018	ISSTA
Power states	Pyles <i>et al.</i> [149]	2011	UbiComp
	Kim <i>et al.</i> [194]	2012	ICOIN
	Ding <i>et al.</i> [195]	2013	SIGMETRICS
	Metri <i>et al.</i> [141]	2014	UbiComp
	Bokhari and Wagner [196]	2016	GECCO
	Rao <i>et al.</i> [197]	2017	HPCA
Displays	Dong <i>et al.</i> [198]	2009	DAC
	Anand <i>et al.</i> [199]	2011	MobiSys
	Lin <i>et al.</i> [200]	2012	TC
	Lin <i>et al.</i> [201]	2014	DAC
	Chen <i>et al.</i> [202]	2014	Computers & graphics
	Huang <i>et al.</i> [203]	2014	ISLPED
	Chen <i>et al.</i> [204]	2014	HotPower
	Nixon <i>et al.</i> [205]	2014	HotPower
	Li <i>et al.</i> [187]	2014	ICSE
	He <i>et al.</i> [206]	2015	Mobicom
	Lin <i>et al.</i> [207]	2017	ISLPED

TABLE 6  
(Continued)

Category	Authors [Ref]	Year	Venue
CPU	Lee <i>et al.</i> [208]	2018	ITMCCJ
	Chang <i>et al.</i> [209]	2019	DAC
	Lin <i>et al.</i> [210]	2019	DAC
	Nagata <i>et al.</i> [98]	2012	UIC
	Bezerra <i>et al.</i> [211]	2013	PM2HW2N
	Chang <i>et al.</i> [212]	2013	TECS
	Tseng <i>et al.</i> [213]	2014	DAC
	Hsiu <i>et al.</i> [214]	2016	TECS
	Li and Mishra [215]	2016	J PARALLEL DISTR COM
	Muhuri <i>et al.</i> [216]	2019	IEEE Trans. Fuzzy Syst.
	Han and Lee [217]	2020	IEEE Access
APIs	Paek <i>et al.</i> [218]	2010	MobiSys
	Zhuang <i>et al.</i> [219]	2010	MobiSys
	Chon <i>et al.</i> [220]	2011	SenSys
	Oshin <i>et al.</i> [221]	2012	TrustCom
	Zhang <i>et al.</i> [222]	2013	IEEE Sensors
	Linares <i>et al.</i> [223]	2014	MSR
Protocols	Ra <i>et al.</i> [224]	2010	MobiSys
	Nurminen [225]	2010	CCNC
	Pyles <i>et al.</i> [226]	2012	UbiComp
	Lee <i>et al.</i> [227]	2012	IEEE Transactions
	Cheng and Hsiu [228]	2013	INFOCOM
	Siekinen <i>et al.</i> [229]	2013	MoVid
	Li <i>et al.</i> [230]	2016	ICSE
System strategies	Chen <i>et al.</i> [231]	2015	Mobicom
	Martins <i>et al.</i> [232]	2015	ATC
Memory management	Duan <i>et al.</i> [39]	2011	IGCC
	Nguyen <i>et al.</i> [233]	2013	UbiComp
	Hussein <i>et al.</i> [234]	2015	Systor
	Zhong <i>et al.</i> [235]	2015	ITCSDI
Programming languages	Nagata <i>et al.</i> [98]	2012	UIC
	Saborido <i>et al.</i> [44]	2018	EMSE

*Prefetching*, e.g., the caching of data transmissions and advertisements in advance, can be used to reduce energy consumption. Balasubramanian *et al.* [175] distinguished applications in delay-tolerant and applications that can benefit from prefetching, to decide which networking technology (3G, GSM, WiFi) to use. Mohan *et al.* [177] and Chen *et al.* [176] prefetched multiple ads to reduce energy consumption induced by downloads. Dutta and Vandermeer [179] achieved energy reductions with caching of up to 45 percent, even with small cache sizes (e.g., 250MB). Yang and Cao [178] formalized the prefetching for energy reductions as an optimization problem. Two approaches (greedy and discrete) are investigated to minimize energy consumption with regard to the network condition (LTE).

*Antipatterns* can be defects such as energy bugs that will likely drain energy. Pathak *et al.* [180] defined energy bugs as errors that cause the system to unexpectedly consume energy. Banerjee *et al.* [147] categorized energy inefficiencies into two categories: energy hotspots and energy bugs.

Energy hotspots cause high battery consumption even though the hardware utilization is low. Energy bugs prevent the idle state of smartphones causing undesired battery consumption without user activity. Pathak *et al.* [180] categorized energy bugs caused by hardware (faulty battery, hardware damage) and software (OS, configurations, applications) and proposed a framework to detect the causes of energy bugs. Pathak *et al.* [181] focused on detecting a particular type of energy bug (no-sleep bug) via static analysis. A no-sleep bug occurs when application components are being kept active when a smartphone is in an idle state, without the necessity of being kept active. Banerjee *et al.* [147] created a framework that automatically generates tests to detect energy bugs. Each test contains a sequence of user interactions that are aimed at revealing energy bugs. As system calls are a primary source for energy bugs, a directed search is used to generate test cases containing system calls. Zhang *et al.* [140] developed ADEL (Automatic Detector of Energy Leaks), to identify energy leaks caused by network operations. Liu *et al.* [182] created GREENDROID, a tool that extends Java PathFinder (JPF) to automatically detect energy problems and report actionable information to combat these problems. Jabbarvand and Malek [183] proposed  $\mu$ DROID, a mutation testing framework, that can be used to detect energy inefficiencies. This framework uses 50 different mutation operators and the similarity of power traces between original application and mutants is used as the test oracle. The detection and removal of energy bugs is not simple, as high energy consumption in applications is not necessarily a sign for wasted computations [140].

*Refactoring*, for example, by using energy-efficient algorithms, can be used to reduce energy consumption. Pathak *et al.* [184] manually restructured the source code of applications to make efficient use of high power states of components. They observed that applications consume I/O energy in distinct lumps. Bundling these lumps can reduce energy consumption. Similarly, Alam *et al.* [186] optimized the placement of wakelock calls. Lyu *et al.* [46] refactored database operations to avoid inefficiencies and reduce energy consumption. Another approach is to change the choice of colors used in an application, as the power consumption of displays is effected by the displayed color [236]. This goes as far that some applications consume double the energy as they would do if colors were optimized for energy consumption [156]. Li *et al.* [187] proposed an approach to automatically change the colors used in web applications. Linares *et al.* [155], [237] used multi-objective optimization to reduce the energy consumption of GUIs, while offering visually similar colors to the original design. Bruce *et al.* [188] applied Genetic Improvement (GI) to find a more energy efficient version of applications. Mutation operations were applied to the source code of a Boolean satisfiability solver, to reduce energy consumption as a measure of fitness. Bokhari *et al.* [164] applied approximate computing on Rebound,<sup>11</sup> a Java Physics library, to achieve a trade-off between accuracy and energy consumption.

Another approach to automatically refactor applications is to follow energy efficiency guidelines [189], [190], [193].

Cito *et al.* [145] adapted application binaries to adjust the frequency of network requests to advertisements and analytics based on the battery status. Anwer *et al.* [185] adapted permissions and corresponding source code of applications based on user requirements, which can for example prevent the unconscious sending of an SMS.

Morales *et al.* [142] showed that there is a correlation between anti-patterns and energy consumption of mobile applications, and proposed the use of multi-objective search to find a set of refactoring sequences able to simultaneously improve code design quality (including the removal of energy smells) and reduce energy consumption. Banerjee *et al.* [191] performed an automatic repair of energy bugs with static and dynamic analysis. Cruz and Abreu [192] manually fixed antipatterns based on Android performance-based guidelines.

*Power states* determine the operating modes of hardware components, which require different amount of energy [153]. Power state transitions can be initiated by hardware components, but are usually performed by the OS [39]. As idle power consumption accounts for approximately 50 percent of the total energy consumption in a smartphone, it is suggested that using different power modes to shut down components is useful to reduce energy consumption [139]. Metri *et al.* [141] developed BATTERYEXTENDER, a tool that enables users to reconfigure device resources to reduce battery consumption. For this purpose, battery consumption of components is predicted with little computational overhead, by using energy profiling. Users are able to pick a period of time for which they want to reduce energy consumption and then choose which components to put in an idle power state to save energy. Bokhari and Wagner [196] proposed a framework to optimize default settings of smartphone components to reduce energy consumption. This problem is formulated as an optimization problem, to minimize energy consumption by changing settings of components based on user behavior. Rao *et al.* [197] dynamically selected system configurations (CPU frequency and memory bandwidth) that reduce energy consumption while maintaining a user-specified level of responsiveness. Ding *et al.* [195] determined power modes based on wireless signal strength, as a poor signal strength drains energy. Based on this, network traffic can be delayed under poor signal strength and continued when the network strength improves. Kim *et al.* [194] limited data transmissions to smartphones (e.g., text, image, and video) based on battery status. By minimizing the amount of transferred data in Social Networking Services, energy consumption can be reduced. For instance, Pyles *et al.* [149] switched Wi-Fi to a low power or sleep mode during periods where it is not being used.

*Displays* are under constant energy consumption while mobile devices are used. Dong *et al.* [198] were the first to study the transformation of GUI colors of OLED displays to reduce energy consumption. Their automatic transformation can be applied on GUI elements (structured) or on pixel information (unstructured). Anand *et al.* [199] adjusted the brightness of screens to reduce the backlight level of the display. Lin *et al.* [200] reduced backlight energy consumption for mobile streaming applications, while other work dimmed areas of the screen [187], [202], [204], [207]. Other

11. <https://github.com/facebookarchive/rebound>

approaches include the adaptation of pixels [201], reduction of frame refreshes [203], [205], [208], [209], pixel density [205], [206] and resolution [210].

CPU clock frequency impacts energy consumption [139]. Nagata *et al.* [98] proposed a method that adjusts CPU clock frequency based on application requirements. Hsiu *et al.* [214] allocated computing resources based on the sensitivity of different applications. Application sensitivity states can be HIGH (interactive), MEDIUM (foreground) or LOW (background). Tseng *et al.* [213] adapted the allocation of CPU resources to applications based on their delay-sensitivity. Further approaches adjust the frequency and voltage of devices (Dynamic Voltage and Frequency Scaling) [211], [212], CPU frequency [217] or the number of cores [215]. Muhuri *et al.* [216] considered linguistic feedback from users to adapt CPU frequency accordingly. They proposed the approach Per-C for Personalized Power Management Approach (Per-C PPMA), which collects user feedback about their degree of satisfaction when using an application. This can be applied to not only reduce energy consumption, but also to improve user satisfaction.

APIs impact energy consumption, as Li *et al.* [238] showed that 91.4 percent of applications consume more than 60 percent of their energy with APIs. Linares *et al.* [223] analyzed usage patterns of “energy-greedy” APIs and give recipes to reduce energy consumption. To support their quantitative and qualitative exploration of API usage pattern, they mined thousands of method calls and API usage patterns. Among those, there are usage patterns that have an unavoidable, high energy consumption, and others which can be improved. An example for energy-greedy APIs is GPS. Paek *et al.* [218] adapted the rate of GPS, and only turns on GPS, when the current location estimate is uncertain. Turning on GPS indoors is also avoided. Other approaches reduce the sampling rate of GPS [219], [220], [221], [222].

Protocols can be used by mobile devices to optimize the energy consumption of networking technologies. Ra *et al.* [224] designed an algorithm to optimize the energy-delay trade-off of delay-tolerant applications that can benefit from low-energy wifi connections. Energy can be reduced if mobile traffic is delayed to a situation where wifi is available [227]. Pyles *et al.* [226] saved energy by prioritizing wifi traffic based on application priority. Li *et al.* [230] bundled HTTP requests to reduce energy consumption. Cheng and Hsiu [228] considered signal strength to reduce energy consumption when fetching location-based information. Nurminen [225] showed that parallel TCP downloading can be used to reduce energy consumption. Siekkinen *et al.* [229] reduced energy consumption of streaming applications by shaping LTE traffic into bursts. Hoque *et al.* [239] surveyed other approaches for optimizing the energy efficiency of streaming.

System strategies can be used to manage background processes. Martins *et al.* [232] introduced TAMER, an OS mechanism that allows rate-limiting of background processes to reduce energy consumption. TAMER imposes on events and signals that cause background applications to wakeup and thereby consume a higher amount of energy. Among others, TAMER can limit the frequency of notifications an application sends while it runs in the background. Chen

*et al.* [231] avoided running applications in the background when they are not beneficial for user experience.

Memory management strategies can be used to change or increase memory to cope with higher requirements of applications. However, a larger main memory size leads to higher energy consumption [39], [131]. Energy reductions can be achieved by using non-volatile memory [235], Phase Change Memory [39] or adaptations to the garbage collection [234] and scheduling algorithms [233].

Programming Languages impact the responsiveness and energy consumption of applications. Nagata *et al.* [98] compared applications developed in different programming languages (Java, JNI and C) and showed that the energy consumption for JNI and C is smaller than for Java. A programming language construct that impacts energy consumption refers to maps (e.g., using HashMap over ArrayMap can reduce energy consumption by 16 percent [44]).

### 7.3 Summary

Energy consumption is a crucial characteristic of embedded systems since these devices have a limited battery size. Energy is consumed by applications (often by multiple applications at the same time), which use several components (e.g., GPS, audio, wifi, display) [24], [39].

To profile energy consumption, two techniques have been pursued. On one hand, energy has been measured with either internal or external instruments [144]. On the other hand, energy consumption has been empirically estimated. For this purpose, various indicators have been investigated [153], [155], [160]. However, when profiling energy consumption, one should consider noise, which impacts the validity of measurements [164], [165], [166].

Many approaches have been proposed to optimize energy consumption. The majority of these approaches address the application source code (38/85). Changes to source code included the removal of bad coding practices (antipatterns) [180], [181], [191], or refactoring to include best practices [189], [190], [193]. Source code adaptation approaches can also focus on particular elements of devices, such as the display. For instance, energy consumption of displays has been reduced by changing colors used in applications [155], [236], [237]. Adaptions have been also applied to displays directly. Thereby, display energy consumption has been reduced by changing brightness, colors and dimming [187], [198], [199], [202], [204], [207].

Similar to responsiveness optimization, energy inefficient computations have been offloaded to external devices with fewer computational restrictions [148]. Adapting the CPU clock frequency, to adjust computational speed, have also reduced energy consumption [98]. An adjustment of components can be performed with different power states (e.g., setting components in an idle state when they are not required).

## 8 DISCUSSION

This section provides an overview and discussion of the optimization approaches presented in previous sections (Sections 4, 5, 6, and 7). In the following, we present a timeline of important methods and techniques proposed for advancing mobile applications' performance. Additionally,



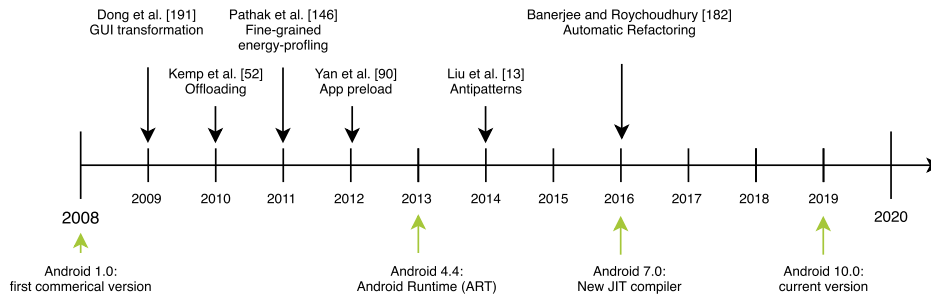


Fig. 4. Timeline of key ideas in non-functional performance optimization and changes to the Android os.

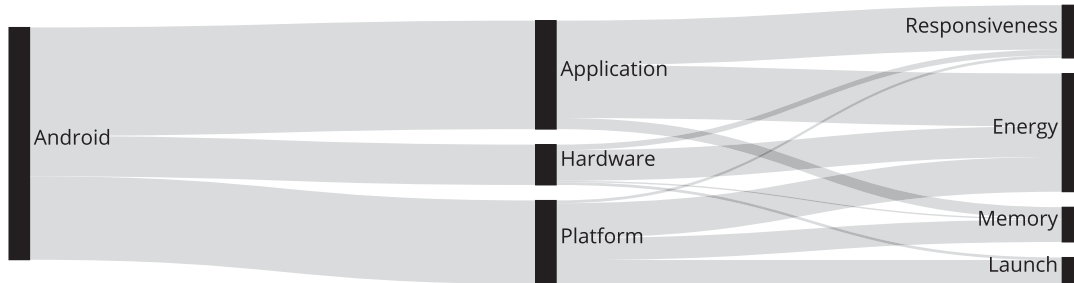


Fig. 5. Flow diagram of optimization approaches. Optimization approaches are matched with the layer of the embedded system (e.g., hardware, platform, and application). These layers are matched with corresponding, optimized non-functional performance characteristics (e.g., responsiveness, launch time, memory, and energy consumption). The width of the flows represents the number of publications with regards to a relationship.

we discuss current challenges and opportunities in this field.

## 8.1 Timeline

Important changes to the Android Platform and publications concerned with optimization of non-functional performance characteristics are shown in the timeline in Fig. 4. The first commercial Android version was introduced in 2008. In 2009, Dong *et al.* [198] studied the transformation of the GUI to reduce the energy consumption of mobile device displays.

In 2010, Kemp *et al.* [66] were the first that proposed an offloading framework for Android applications. This framework can be used to offload computationally expensive parts of mobile applications and improve both responsiveness and energy consumption.

In 2011, Pathak *et al.* [153] extended energy profiling of hardware utilization based on system-calls to provide fine-grained energy estimates. In 2012, Yan *et al.* [101] proposed an OS extension that preloads applications and application-specific content to improve the launch time of applications. In 2013, Android designers applied a big change to the Android platform, introducing ART over the previously used DVM. With this change, optimizations such as Ahead-Of-Time (AOT) compilation are applied by the Android platform to improve application performance.

In 2014, Liu *et al.* [13] characterized several types of performance bugs, which addressed different non-functional characteristics of Android applications and caused excessive resource consumption of memory and battery. Examples of performance bugs include GUI lagging, energy leaks, and memory bloat. Based on the knowledge of bugs and antipatterns or guidelines, Banerjee and Roychoudhury [189]

proposed the automated refactoring of application source code in 2016.

In 2016, Android 7.0-7.1 introduced a new JIT Compiler. This allowed faster application installations and it reduced the size of compiled code. The 10th major Android version was published in 2019.

## 8.2 Optimization Approaches per Android Layer

Reading the related work, we observed that approaches are applied to different layers of mobile devices. The distribution of these approaches differs between the four non-functional performance characteristics we investigated. In Fig. 5, we organize optimization approaches based on layers (application, platform, hardware) to non-functional performance characteristics (i.e., responsiveness, launch time, memory and energy).

The majority of approaches to optimize responsiveness are application-based techniques and apply changes to the source code (e.g., antipatterns, refactoring or prefetching). Hardware and I/O optimization approaches are concerned with increasing the speed of reads and writes, faster fetching of bytecode as well as correct usages of CPU and GPU.

Launch time has not been improved itself by changes to applications, but, in majority, by changes in the Android platform. In particular, the preloading of applications [16], [37], [47], [48], [100], [101], [106], [107], [108], [240] and better choices for the LMK [15], [37], [47], [48], [103], [110], [111], [112], [113], to prevent cold starts from happening, are investigated frequently. The overall research direction is concerned with reducing the amount of cold starts rather than reducing the launch time itself.

Memory is mostly optimized by removing antipatterns [50], [51], [121], [123], [124] and applying different

strategies for the garbage collection [38], [51], [95], [109], [125], [126], [127].

A great deal of research work optimizes energy consumption and almost every optimization category has been addressed for energy consumption. Unlike the optimization of responsiveness, launch time and memory usage, a large portion of approaches apply changes to the hardware, especially for screens. The power mode of components plays an important role as well, which is governed by the Android OS. Furthermore, API usage has a huge impact on energy consumption. In particular, the use of GPS. Alike responsiveness, changes to application source code contribute to energy savings. Approaches like offloading optimize both, responsiveness and energy consumption.

### 8.3 Relationship of Optimization Approaches

A major trade-off that can be seen is between energy consumption and responsiveness [197], [215], [241], which should therefore be evaluated in union. One example for the relationship of responsiveness and energy consumption can be seen in the CPU clock frequency, as reducing it leads to less energy consumption, but lowers responsiveness and vice versa [98].

Considering application launch, hot starts are preferable over cold starts, not only to reduce the launch time but also to reduce energy consumption [25], [37], [242].

The choice of the memory system and storage technique used impacts energy consumption [39], [243]. One aspect of this is the size of memory, as a larger main memory size leads to higher energy consumption [39], [131]. Regardless, larger memory permits the preloading and prefetching of more application data, which can improve launch time, responsiveness and/or energy consumption. Evidence for potential improvement comes from the observation that the number of applications cached directly impacts the amount of blocked memory [110]. Memory leaks [242] or a full memory [15] limit the number of preloaded applications. Different memory management strategies for the LMK [47], [48] have shown to improve application launch.

While there is a relationship among non-functional characteristics, several optimization approaches improve more than one characteristic at the same time, which we will present in the following.

Gordon *et al.* [71], Chun *et al.* [67], Chen and Hao [73], Kemp *et al.* [66] and Kosta *et al.* [69] performed offloading to improve responsiveness, and as a result reduced energy consumption. Khairy *et al.* [170], and Qian and Andresen [174] proposed offloading systems aimed at reducing energy consumption, that reduced execution time as well. In order to achieve improvements with offloading, the saving has to exceed the additional cost imposed by the offloading process [244].

Removing antipatterns can reduce the number of garbage collection calls as well [50], whereas memory leaks reduce available memory and responsiveness [124].

Nguyen *et al.* [49] improved both responsiveness and launch time by optimizing I/O read and write operations. Kim *et al.* [132] reduced CPU usage and energy consumption with a region-based physical memory management scheme. Hecht *et al.* [50] improved memory as well as user interface

performance by correcting code smells. The approach of Han *et al.* [217] on CPU scheduling could not only reduce energy consumption but also application launch time. Hsiu *et al.* [214] reduced application response time by scheduling computing resources for energy reduction. Saborido *et al.* [44] showed that the choice of different map implementations impacts memory, energy and responsiveness. Lyu *et al.* [46] refactored database operations and thereby improved responsiveness and energy consumption. By reducing buffer cache, Kim and Bahn [134] reduced energy consumption at the same time.

Optimization approaches for responsiveness, launch time and memory can induce additional energy costs, if they extend the OS or require additional computations [28], [84], [94], [101], [106], [108], [128]. Overhead is not only produced by optimization approaches, but also by profiling and testing [59], [141], [159], [245], as well as by displaying ads [246] to generate revenue. In addition to causing overhead, non-functional improvements may come with a trade-off with regards to functional characteristics, such as the accuracy of algorithms [150].

### 8.4 Challenges

Reflecting on the different optimization approaches and the relationship between non-functional performance characteristics, we have identified challenges and opportunities for future work. We begin by outlining three challenges that developers face when optimizing non-functional performance characteristics, followed by an overview of future work targeted by the surveyed publications, and by opportunities we detected.

*Cross-Characteristic Dependencies.* Section 8.3 shows that a challenge that developers face while optimizing applications' performance refers to the handling of the dependencies among different performance characteristics. This means that while improving one characteristic, they may decrease the performance of another [39], [98], [131]. For example, optimization approaches that extend the OS or invoke additional computation cause additional energy consumption [28], [84], [94], [101], [106], [108], [128]. Therefore, there should be a balance that developers need to achieve among performance characteristics. While we pointed out works that achieved improvements in more than one performance characteristic (Section 8.3), it would be interesting if adverse relationships and trade-offs between non-functional performance characteristics receive more attention. One example for this is the energy consumption and responsiveness trade-off when adjusting CPU clock frequency [98].

*Testing Cost.* There are different definitions of the four non-functional characteristics to determine how performance is measured (e.g., responsiveness measured in ms [55], [56] or frames [50], or different definitions of application launch completion [100], [101]). Another problem arises with noisy measurements (e.g., due to hardware components [164]), as we have seen in Section 7. Therefore, testing should consider variance in measurements to ensure statistical significance [165], [166].

An attempt to reduce the cost of testing is the usage of emulators [122], [163] or prediction of non-functional performance without executing an application. Prediction of

performance has been applied for responsiveness [56], [65] and energy consumption [7], [153], [154], [155], [156], [157], [158], [159], [160].

*User Satisfaction.* While improvements in non-functional characteristics without performance deterioration in other characteristics, can always be seen as something positive, quantifying the impact of improvements on user satisfaction remains challenging.

Two of the surveyed publications attempted to tackle this issues. For example, Muhuri *et al.* [216] collected linguistic feedback about user satisfaction (e.g., ranging from “very low” to “extremely high”) to examine the impact of performance on user satisfaction. This information has been used to adapt the CPU frequency. Zhao *et al.* [105] stated that collecting user feedback can be costly and inconvenient for developers. To overcome this issue, they mapped a user-perceived satisfaction score about launch times to the actual launch time delay, which is easily measurable. Such an incorporation of user satisfaction in the optimization procedure, could prove to be an interesting consideration for future work on other performance characteristics.

In addition to these challenges, the reviewed publications distinguished several fields of future work, including:

- Improvement and extension of prediction methods (e.g., for offloading, prefetching, and preloading) [37], [65], [66], [85], [114], [115], [131], [220];
- Investigation of antipatterns [13], [50], [57], [122], [223];
- Automation of the optimization process [50], [147]
- Extension of testing (e.g., usage of more devices and applications) [50], [182], [206];
- Improvement of measurements (profiling) [131], [188].

In particular, extensions of prediction methods include the consideration of additional information, such as the context (e.g., location, time) [37], behavioral patterns [85] and information about remote resources for offloading (e.g., processor speed, available memory) [66]. Future work on antipatterns is concerned with investigating a broader range of antipatterns [50], [57] and discovering new antipatterns or categories [13], [122], [223]. Automation could be applied to time-consuming tasks, such as finding causes for energy wastage [147] or the correction of antipatterns [147].

Based on our results, including Fig. 5, we identified further gaps in the literature. For once, changes to applications are not investigated for launch time improvements. It could be investigated whether antipatterns, that exist for responsiveness, memory and energy usage, exist for application launch as well.

Automatic refactoring has been applied for both responsiveness and energy, individually. It could therefore be interesting to apply refactoring in a multi-objective setting, to optimize both. Lastly, changes to the platform have scarcely been used to improve responsiveness. There could be the potential to apply ideas from other non-functional characteristics, such as APIs for energy consumption.

## 9 RELATED WORK

In the following, we give an overview of literature related to the non-functional performance optimization of Android

applications. At first, we look into optimization approaches for software engineering. We furthermore describe the developer and user perspective on mobile-application optimization.

### 9.1 Optimization in Software Engineering

This survey outlined optimization approaches for mobile applications. In the following, we present studies and insights on optimization with regards to software engineering.

During the development and optimization of software, it is important to consider software requirements [247], which can contain functional, non-functional, business and user requirements. Different techniques for prioritizing requirements [248], [249] can be applied to determine the importance of non-functional over functional characteristics.

Nonetheless, the performance of software is difficult to be measured, as it is pervasive and affected by various different aspects (e.g., the platform used) [250]. Software Performance Engineering (SPE) is an approach to measure and improve system performance [250]. SPE subsumes software engineering activities that are applied to meet performance requirements and achieve improvements. Profiling tools can be used to measure performance (e.g., for energy consumption [251], [252], GPUS [253], responsiveness [254], and memory [255]).

Building upon performance profilers, optimization and improvement techniques can be applied to software. Petke *et al.* [32] conducted a survey on genetic improvement of software. They mention improvements for non-functional characteristics for energy and memory consumption as well as functional improvements including repairs and addition of new functionalities.

Following, we outline representative publications on software optimization. Mao and Humphrey [256] analyzed long, unexpected startup times of Virtual Machines on the cloud. Kaminaga [257] discussed techniques to achieve a faster startup time for embedded Linux systems. Van Emden and Moonen [258] analyzed code smells in Java source code. Baer and Chen [259] proposed a hardware scheme for preloading data. Sahin *et al.* [260] analyzed the impact of code refactoring on the energy consumption of 197 applications. Chen *et al.* [261] used flash drives in mobile computers for caching and prefetching data to save energy. Chen *et al.* [262] reduced writes to memory by using deduplication. Hauswirth and Chilimbi [263] detected memory leaks with low computational overhead. Padmanabhan and Mogul [264] investigated latency reductions by prefetching web documents.

### 9.2 Developer Perspective

In the following, we outline the developer perspective on optimization and development of mobile applications, including challenges during the development process. While certain characteristics of smartphones and PCs are similar [104], developers encounter differences to conventional software development as mobile applications are smaller than traditional software [265].



Developers use tools to support the development of applications [266], profiling, and debugging [267]. Static analysis can be used to support developers to find bugs and inspect code [268]. Other tools perform security assessment, automated test case generation and detection of non-functional issues such as energy consumption [269], [270]. While fixing non-functional performance bugs, developers need to consider the threat of introducing functional bugs [271] and hindering code maintainability [272]. In this context, Linares *et al.* [273] suggested that developers rarely implement micro-optimizations (e.g., changes at statement level).

When developing an application, developers need to decide which and how many platforms to use (e.g., Android, ios, Windows os). Note that each platform faces non-functional issues alike (e.g., antipatterns can be found in ios [274], [275]). This impacts the development effort, as multiple codebases need to be maintained. Development of applications for multiple platforms can be supported by cross-platform tools (CPTS) [276]. By doing so, developers compromise between user experience and the ability to publish an application on multiple platforms. Willocx *et al.* [54] found that CPTS lead to an increased launch time of applications.

Furthermore, devices vary in their available memory, CPU, and display size, which has implications on application performance [277]. Therefore, developers need to test their applications on multiple devices.

Further changes and added functionality to the Android OS can be imposed by phone vendors [278]. This leads to varying behavior across different smartphone types [279]. Khalid *et al.* [280] analyzed the number of different phones that use applications in order to help developers to decide how many devices to use when testing applications, as the rating varies for different phone types. Often, applications run on more than hundred different phone types, implying a huge computational effort if all devices would be tested for. They found that around a third of the devices account for 80 percent of the reviews and thereby usage, which can be used to prioritize which devices to use during testing. Lu *et al.* [281] prioritized devices to test applications based on the amount of user activity rather than the number of devices.

In order to analyze the performance of applications, the Google Play Store provides developers with pre-launch reports after an application is published. Tests are carried out on different devices and for up to five languages [282]. Another tool that developers can use to judge the quality and performance of applications are Android Vitals. Statistics including battery consumption, and crashes are collected from real users and reported to developers [29]. Comparisons with regard to non-functional performance characteristics, such as energy consumption, can also be measured with applications of the same category [283].

Frequently, developers use user reviews and test applications manually to detect performance issues and bugs [267]. Developers can change applications based on user reviews. Those reviews contain among others information about performance, bugs, problems or new features [3], [30]. There is a huge number of reviews that are written for applications, where some of them contain relevant information for developers [284]. For this purpose, Chen *et al.* [19] developed a

framework to filter informative reviews by applying text mining and ranking methods. Furthermore, reviews can be analyzed for trends [285], [286] and emerging issues [287].

Even though developer identities do not often impact the choice of applications (only 11 percent of users choose an application based on who developed it [2]), they significantly impact the quality and success of applications [1]. Other factors that are correlated with the rating of applications include the apk size, minimum required SDK, and number of images on the application description page [288].

### 9.3 User Perspective

In the context of mobile applications, user shows a high degree of individualism [289]. Therefore, not every optimization approach can be applied in a general fashion. This is why understanding user behavior and differences among users and user groups is important when improving performance. Ultimately it is the user who decides the changes to an application result in a higher level of satisfaction.

The user-perceived quality of an application is not only determined by the application itself, but also by the device and attributes of its components [290]. Aspects that lead to the most negative complaints in reviews are related to privacy, hidden costs and features of the application [3].

Users' behavior varies in terms of number of interactions, amount of data received, interaction length and number of applications used [7]. Some users even show addictive behavior [291]. Approaches that aim to improve user experience (e.g., by reducing energy consumption or improving the responsiveness), should therefore be adaptive to user behavior. Making matters more difficult to predict, usage patterns can change within a few days [116].

Additionally, user behavior across countries shows significant differences as well [2]. Users of different countries prioritize other aspects of applications and show variations in rating applications as well as writing reviews [2]. For example, users in China are more likely to rate an application, while users in Brazil are more likely to abandon a slow or buggy application. While there exist differences among Android users based on countries, a study on application launch performed by Morrison *et al.* [292] showed that similarities between Android and ios exist.

Application usage can be seen as a sequence, with multiple applications used consecutively in a short period of time, whereas 68.2 percent of sequences only contain a single application [114]. Application usage varies based on the context, such as location [293] and/or time [114], [294].

## 10 THREATS TO VALIDITY

In this section, we discuss potential threats to our survey based on internal and external validity.

*Internal validity* refers to problems of our methods that could threaten the validity of results and claims made in this survey [9]. A potential threat to internal validity is the completeness of the reviewed literature. Ideally, every relevant publication is included after the search process; however, the risk of missing a publication cannot be eliminated. A relevant publication can be missed if the respective data source containing the publication is not properly searched. Another cause for missing a publication is that

corresponding keywords are not included in our searching criteria. To address both causes, we performed a preliminary search to gather relevant keywords and venues to guide the literature search. Moreover, two authors independently carried out the filtering process and their results were cross-checked, in order to ensure reliability and reduce researchers' bias. A different threat to internal validity refers to the precision of our results (e.g., the inclusion of irrelevant publications). To mitigate the impact of irrelevant publications, we check the title, abstract, and body of the publications examined as outlined in Section 3.1.3. Furthermore, there is a risk of drawing incorrect conclusions or claims. For this purpose, every stage of the search and analysis of results (e.g., repository search and categorization of approaches) has been performed by one author and cross-checked by another.

External validity describes the generalizability of our results outside of the given scope [295]. A potential, external threat is that the chosen non-functional performance characteristics are insufficient to describe optimization techniques for embedded systems. Through our preliminary search and analysis of related work (Section 9) we found that the selected four non-functional characteristics (responsiveness, launch time, memory and energy consumption) are representative. Another threat is the applicability of our study to other mobile platforms (e.g., ios). While there are differences between ios and Android, the general organization and functionality within embedded systems remain the same. We therefore argue that our categorization can be applied to mobile platforms, other than Android.

## 11 CONCLUSION

In this paper, we have provided an overview of the existing research work on non-functional performance optimization for Android applications published between 2008 and 2020. Our survey presents optimization approaches for non-functional performance characteristics (e.g., responsiveness, launch time, memory, and energy). It also shows relationships among these characteristics, and identifies research gaps for potential future works. We hope that this survey will help researchers and developers to have a holistic perception on optimization approaches for mobile devices, the impact of these approaches, and the significance of different performance characteristics.

## ACKNOWLEDGMENTS

This work was supported by ERC Advanced fellowship under Grant 741278 (EPIC). We would like to thank the members of the community who kindly provided comments and feedback on an earlier draft of this article. Our living surveys[21] has been inspired by Allamanis *et al.* [296].

## REFERENCES

- [1] V. Inukollu, D. Keshamoni, T. Kang, and M. Inukollu, "Factors influencing quality of mobile apps: Role of mobile app development life cycle," *Int. J. Softw. Eng. Appl.*, vol. 5, pp. 15–34, 2014.
- [2] S. L. Lim, P. J. Bentley, N. Kanakam, F. Ishikawa, and S. Honiden, "Investigating country differences in mobile app user behavior and challenges for software engineering," *IEEE Trans. Softw. Eng.*, vol. 41, no. 1, pp. 40–64, Jan. 2015.
- [3] H. Khalid, E. Shihab, M. Nagappan, and A. E. Hassan, "What do mobile app users complain about?," *IEEE Softw.*, vol. 32, no. 3, pp. 70–77, May/Jun. 2015.
- [4] A. AlSubaihini, F. Sarro, S. Black, L. Capra, and M. Harman, "App store effects on software engineering practices," *IEEE Trans. Softw. Eng.*, vol. 47, no. 2, pp. 300–319, Feb. 2021.
- [5] A. Mazuera-Rozo, C. Trubiani, M. Linares-Vásquez, and G. Bavota, "Investigating types and survivability of performance bugs in mobile apps," *Empirical Softw. Eng.*, vol. 25, pp. 1644–1686, 2020.
- [6] A. Finkelstein, M. Harman, Y. Jia, W. Martin, F. Sarro, and Y. Zhang, "Investigating the relationship between price, rating, and popularity in the blackberry world app store," *Inf. Softw. Technol.*, vol. 87, pp. 119–139, 2017.
- [7] H. Falaki, R. Mahajan, S. Kandula, D. Lymberopoulos, R. Govindan, and D. Estrin, "Diversity in smartphone usage," in *Proc. 8th Int. Conf. Mobile Syst., Appl., Serv.*, 2010, pp. 179–194.
- [8] P. Welke, I. Andone, K. Blaszkiewicz, and A. Markowetz, "Differentiating smartphone users by app usage," in *Proc. ACM Int. Joint Conf. Pervasive Ubiquitous Comput.*, 2016, pp. 519–523.
- [9] M. Kechagia, D. Mitropoulos, and D. Spinellis, "Charting the API minefield using software telemetry data," *Empirical Softw. Eng.*, vol. 20, no. 6, pp. 1785–1830, 2015.
- [10] C. Hu and I. Neamtiu, "Automating GUI testing for android applications," in *Proc. 6th Int. Workshop Automat. Softw. Test*, 2011, pp. 77–83.
- [11] P. Bhattacharya, L. Ulanova, I. Neamtiu, and S. C. Koduru, "An empirical analysis of bug reports and bug fixing in open source android apps," in *Proc. 17th Eur. Conf. Softw. Maintenance Reeng.*, 2013, pp. 133–143.
- [12] W. J. Martin, F. Sarro, and M. Harman, "Causal impact analysis for app releases in google play," in *Proc. 24th ACM SIGSOFT Int. Symp. Found. Softw. Eng.*, Seattle, WA, USA, 2016, pp. 435–446.
- [13] Y. Liu, C. Xu, and S.-C. Cheung, "Characterizing and detecting performance bugs for smartphone applications," in *Proc. 36th Int. Conf. Softw. Eng.*, 2014, pp. 1013–1024.
- [14] F. Ferrucci, C. Gravino, P. Salza, and F. Sarro, "Investigating functional and code size measures for mobile applications: A replicated study," in *Proc. Int. Conf. Product-Focused Softw. Process Improvement*, 2015, pp. 271–287.
- [15] H. Kim, H. Lim, D. Manatunga, H. Kim, and G.-H. Park, "Accelerating application start-up with nonvolatile memory in android systems," *IEEE Micro*, vol. 35, no. 1, pp. 15–25, Jan./Feb. 2015.
- [16] A. L. N. Martins, C. A. Duarte, and J. Jeong, "Improving application launch performance in smartphones using recurrent neural network," in *Proc. Int. Conf. Mach. Learn. Technol.*, 2018, pp. 58–62.
- [17] C. Gao, J. Zeng, F. Sarro, M. R. Lyu, and I. King, "Exploring the effects of ad schemes on the performance cost of mobile phones," in *Proc. 1st Int. Workshop Adv. Mobile App Anal.*, 2018, pp. 13–18.
- [18] C. Gao, J. Zeng, F. Sarro, D. Lo, M. R. Lyu, and I. King, "Do users care about ad's performance costs? Exploring the effects of the performance costs of in-app ads on user experience," *Inf. Softw. Technol.*, vol. 132, 2021, Art. no. 106471.
- [19] N. Chen, J. Lin, S. C. Hoi, X. Xiao, and B. Zhang, "AR-miner: Mining informative reviews for developers from mobile app marketplace," in *Proc. 36th Int. Conf. Softw. Eng.*, 2014, pp. 767–778.
- [20] A. Banerjee and A. Roychoudhury, "Future of mobile software for smartphones and drones: Energy and performance," in *Proc. IEEE/ACM 4th Int. Conf. Mobile Softw. Eng. Syst.*, 2017, pp. 1–12.
- [21] M. Hort, M. Kechagia, F. Sarro, and M. Harman, "A survey of performance optimization for mobile applications," [Online]. Available: <https://solar.cs.ucl.ac.uk/os/appoptimization.html>
- [22] R. Cohen and T. Wang, *Overview of Embedded Application Development for Intel Architecture*. Berkeley, CA, USA: Apress, 2014, pp. 1–18.
- [23] N. D. Lane, E. Miluzzo, H. Lu, D. Peebles, T. Choudhury, and A. T. Campbell, "A survey of mobile phone sensing," *Comm. Mag.*, vol. 48, no. 9, pp. 140–150, 2010.
- [24] L. Zhang, B. Tiwana, Z. Qian, Z. Wang, R. P. Dick, Z. M. Mao, and L. Yang, "Accurate online power estimation and automatic battery behavior based power model generation for smartphones," in *Proc. 8th IEEE/ACM/IFIP Int. Conf. Hardware/Softw. Codes. Syst. Synth.*, 2010, pp. 105–114.

- [25] S. Kwon, S.-H. Kim, J.-S. Kim, and J. Jeong, "Managing GPU buffers for caching more apps in mobile systems," in *Proc. 12th Int. Conf. Embedded Softw.*, 2015, pp. 207–216.
- [26] A. Arcangeli, I. Eidus, and C. Wright, "Increasing memory density by using KSM," in *Proc. Linux Symp.*, 2009, pp. 19–28.
- [27] N. Gupta, "Compcache: In-memory compressed swapping," 2009. [Online]. Available: <https://lwn.net/Articles/334649/>
- [28] B. Lee, S. M. Kim, E. Park, and D. Han, "Memscope: Analyzing memory duplication on android systems," in *Proc. 6th Asia-Pacific Workshop Syst.*, 2015, p. 19.
- [29] J. Harty and M. Müller, "Better android apps using android vitals," in *Proc. 3rd ACM SIGSOFT Int. Workshop App Market Anal.*, 2019, pp. 26–32.
- [30] A. Di Sorbo *et al.*, "What would users change in my app? Summarizing app reviews for recommending software changes," in *Proc. 24th ACM SIGSOFT Int. Symp. Found. Softw. Eng.*, 2016, pp. 499–510.
- [31] M. Harman, Y. Jia, and Y. Zhang, "App store mining and analysis: MSR for app stores," in *Proc. 9th IEEE Work. Conf. Mining Softw. Repositories*, 2012, pp. 108–111.
- [32] J. Petke, S. O. Haraldsson, M. Harman, W. B. Langdon, D. R. White, and J. R. Woodward, "Genetic improvement of software: A comprehensive survey," *IEEE Trans. Evol. Comput.*, vol. 22, no. 3, pp. 415–432, Jun. 2018.
- [33] B. W. Boehm, J. R. Brown, and Hans Kaspar, *Characteristics of Software Quality*. Amsterdam, The Netherlands: North-Holland, 1978, vol. 1.
- [34] G.-C. Roman, "A taxonomy of current issues in requirements engineering," *Computer*, vol. 18, no. 4, pp. 14–23, 1985.
- [35] R. B. Grady and D. L. Caswell, *Software Metrics: Establishing a Company-Wide Program*. Upper Saddle River, NJ, USA: Prentice-Hall, 1987.
- [36] L. Chung and J. C. S. do Prado Leite, "On non-functional requirements in software engineering," in *Conceptual Modeling: Foundations and Applications*. Berlin, Germany: Springer, 2009, pp. 363–379.
- [37] W. Song, Y. Kim, H. Kim, J. Lim, and J. Kim, "Personalized optimization for android smartphones," *ACM Trans. Embedded Comput. Syst.*, vol. 13, no. 2s, 2014, Art. no. 60.
- [38] Y. He, C. Yang, and X.-F. Li, "Improve google android user experience with regional garbage collection," in *IFIP Int. Conf. Netw. Parallel Comput.*, 2011, pp. 350–365.
- [39] R. Duan, M. Bi, and C. Gniady, "Exploring memory energy optimizations in smartphones," in *Proc. Int. Green Comput. Conf. Workshops*, 2011, pp. 1–8.
- [40] B. Kitchenham, "Procedures for performing systematic reviews," *Dept. Comput. Sci., Keele Univ., Keele, U.K., Tech. Rep. TR/SE-0401*, 2004, vol. 33, no. 2004, pp. 1–26.
- [41] A. Sadeghi, H. Bagheri, J. Garcia, and S. Malek, "A taxonomy and qualitative comparison of program analysis techniques for security assessment of android software," *IEEE Trans. Softw. Eng.*, vol. 43, no. 6, pp. 492–530, Jun. 2017.
- [42] W. Martin, F. Sarro, Y. Jia, Y. Zhang, and M. Harman, "A survey of app store analysis for software engineering," *IEEE Trans. Softw. Eng.*, vol. 43, no. 9, pp. 817–847, Sep. 2017.
- [43] C. Wohlin, "Guidelines for snowballing in systematic literature studies and a replication in software engineering," in *Proc. 18th Int. Conf. Evaluation Assessment Softw. Eng.*, 2014, pp. 1–10.
- [44] R. Saborido, R. Morales, F. Khomh, Y.-G. Guéhéneuc, and G. Antoniol, "Getting the most from map data structures in android," *Empirical Softw. Eng.*, vol. 23, no. 5, pp. 2829–2864, 2018.
- [45] P. K. Das, S. Shome, and A. K. Sarkar, "Apps: Accelerating performance and power saving in smartphones using code offload," in *Proc. IEEE 6th Int. Conf. Adv. Comput.*, 2016, pp. 759–765.
- [46] Y. Lyu, D. Li, and W. G. Halfond, "Remove rats from your code: Automated optimization of resource inefficient database writes for mobile applications," in *Proc. 27th ACM SIGSOFT Int. Symp. Softw. Testing Anal.*, 2018, pp. 310–321.
- [47] Y.-F. Chung, Y.-T. Lo, and C.-T. King, "Enhancing user experiences by exploiting energy and launch delay trade-off of mobile multimedia applications," *ACM Trans. Embedded Comput. Syst.*, vol. 12, no. 1s, pp. 1–19, 2013.
- [48] J. Lee, K. Lee, E. Jeong, J. Jo, and N. B. Shroff, "CAS: Context-aware background application scheduling in interactive mobile systems," *IEEE J. Sel. Areas Commun.*, vol. 35, no. 5, pp. 1013–1029, May 2017.
- [49] D. T. Nguyen *et al.*, "Reducing smartphone application delay through read/write isolation," in *Proc. Annu. Int. Conf. Mobile Syst., Appl., Serv.*, 2015, pp. 287–300.
- [50] G. Hecht, N. Moha, and R. Rouvoy, "An empirical study of the performance impacts of android code smells," in *Proc. Int. Conf. Mobile Softw. Eng. Syst.*, 2016, pp. 59–69.
- [51] K. Tasneem, A. Siddiqui, and A. Liaquat, "Android memory optimization," *Int. J. Comput. Appl.*, vol. 975, pp. 36–43, 2019.
- [52] S. Yang, D. Yan, and A. Rountev, "Testing for poor responsiveness in android applications," in *Proc. 1st Int. Workshop Eng. Mobile-Enabled Syst.*, 2013, pp. 1–6.
- [53] N. Tolia, D. G. Andersen, and M. Satyanarayanan, "Quantifying interactive user experience on thin clients," *Computer*, vol. 39, no. 3, pp. 46–52, 2006.
- [54] M. Willocx, J. Vossaert, and V. Naessens, "Comparing performance parameters of mobile app development strategies," in *Proc. IEEE/ACM Int. Conf. Mobile Softw. Eng. Syst.*, 2016, pp. 38–47.
- [55] J. Huang, Q. Xu, B. Tiwana, Z. M. Mao, M. Zhang, and P. Bahl, "Anatomizing application performance differences on smartphones," in *Proc. 8th Int. Conf. Mobile Syst., Appl., Serv.*, 2010, pp. 165–178.
- [56] J. Zhang, X. Wang, and Y. Chen, "Android app performance detection framework based on dynamic analysis of function call graphs," in *Proc. World Symp. Softw. Eng.*, 2019, pp. 1–5.
- [57] W. Zhao, Z. Ding, M. Xia, and Z. Qi, "Systematically testing and diagnosing responsiveness for android apps," in *Proc. IEEE Int. Conf. Softw. Maintenance Evol.*, 2019, pp. 449–453.
- [58] M. Cho, H. J. Lee, M. Kim, and S. W. Kim, "Androscope: An insightful performance analyzer for all software layers of the android-based systems," *Electron. Telecommun. Res. Inst. J.*, vol. 35, no. 2, pp. 259–269, 2013.
- [59] L. Ravindranath, J. Padhye, S. Agarwal, R. Mahajan, I. Obermiller, and S. Shayandeh, "Appinsight: Mobile app performance monitoring in the wild," in *Proc. 10th USENIX Symp. Operating Syst. Des. Implementation*, 2012, pp. 107–120.
- [60] D. K. Hong, A. Nikraves, Z. M. Mao, M. Ketkar, and M. Kishinevsky, "Perfprobe: A systematic, cross-layer performance diagnosis framework for mobile platforms," in *Proc. IEEE/ACM 6th Int. Conf. Mobile Softw. Eng. Syst.*, 2019, pp. 50–61.
- [61] H. Kim, B. Choi, and W. E. Wong, "Performance testing of mobile applications at the unit test level," in *Proc. 3rd IEEE Int. Conf. Secure Softw. Integration Rel. Improvement*, 2009, pp. 171–180.
- [62] Y. Kang, Y. Zhou, H. Xu, and M. R. Lyu, "PersisDroid: Android performance diagnosis via anatomizing asynchronous executions," 2015, *arXiv:1512.07950*.
- [63] —, "DiagDroid: Android performance diagnosis via anatomizing asynchronous executions," in *Proc. 24th ACM SIGSOFT Int. Symp. Found. Softw. Eng.*, 2016, pp. 410–421.
- [64] Y. Wang and A. Rountev, "Profiling the responsiveness of android applications via automated resource amplification," in *Proc. IEEE/ACM Int. Conf. Mobile Softw. Eng. Syst.*, 2016, pp. 48–58.
- [65] Y. Kwon *et al.*, "Mantis: Automatic performance prediction for smartphone applications," in *Proc. USENIX Conf. Annu. Tech. Conf.*, 2013, pp. 297–308.
- [66] R. Kemp, N. Palmer, T. Kielmann, and H. Bal, "Cuckoo: A computation offloading framework for smartphones," in *Int. Conf. Mobile Comput., Appl., Serv.*, 2010, pp. 59–79.
- [67] B.-G. Chun, S. Ihm, P. Maniatis, M. Naik, and A. Patti, "CloneCloud: Elastic execution between mobile device and cloud," in *Proc. 6th Conf. Comput. Syst.*, 2011, pp. 301–314.
- [68] M.-R. Ra, A. Sheth, L. Mummert, P. Pillai, D. Wetherall, and R. Govindan, "Odessa: Enabling interactive perception applications on mobile devices," in *Proc. 9th Int. Conf. Mobile Syst., Appl., Serv.*, 2011, pp. 43–56.
- [69] S. Kosta, A. Aucinas, P. Hui, R. Mortier, and X. Zhang, "ThinkAir: Dynamic resource allocation and parallel execution in the cloud for mobile code offloading," in *Proc. IEEE INFOCOM*, 2012, pp. 945–953.
- [70] M. S. Gordon, D. A. Jamshidi, S. Mahlke, Z. M. Mao, and X. Chen, "COMET: Code offload by migrating execution transparently," in *Proc. 10th USENIX Symp. Operating Syst. Des. Implementation*, 2012, pp. 93–106.
- [71] M. S. Gordon, D. K. Hong, P. M. Chen, J. Flinn, S. Mahlke, and Z. M. Mao, "Accelerating mobile applications through flip-flop replication," in *Proc. 13th Annu. Int. Conf. Mobile Syst., Appl., Serv.*, 2015, pp. 137–150.



- [72] R. Montella *et al.*, "Accelerating linux and android applications on low-power devices through remote GPGPU offloading," *Concurrency Comput.: Pract. Experience*, vol. 29, no. 24, 2017, Art. no. e4286.
- [73] M. Chen and Y. Hao, "Task offloading for mobile edge computing in software defined ultra-dense network," *IEEE J. Sel. Areas Commun.*, vol. 36, no. 3, pp. 587–597, Mar. 2018.
- [74] G. Jin, L. Song, X. Shi, J. Scherpelz, and S. Lu, "Understanding and detecting real-world performance bugs," *ACM SIGPLAN Notices*, vol. 47, no. 6, pp. 77–88, 2012.
- [75] A. Nistor, L. Song, D. Marinov, and S. Lu, "Toddler: Detecting performance problems via similar memory-access patterns," in *Proc. Int. Conf. Softw. Eng.*, 2013, pp. 562–571.
- [76] T. Ongkositi and S. Takada, "Responsiveness analysis tool for android application," in *Proc. 2nd Int. Workshop Softw. Develop. Lifecycle Mobile*, 2014, pp. 1–4.
- [77] G. Hecht, O. Benomar, R. Rouvroy, N. Moha, and L. Duchien, "Tracking the software quality of android applications along their evolution (T)," in *Proc. 30th IEEE/ACM Int. Conf. Automated Softw. Eng.*, 2015, pp. 236–247.
- [78] S. Habchi, X. Blanc, and R. Rouvroy, "On adopting linters to deal with performance concerns in android apps," in *Proc. 33rd ACM/IEEE Int. Conf. Automated Softw. Eng.*, 2018, pp. 6–16.
- [79] W. Li, Y. Jiang, C. Xu, Y. Liu, X. Ma, and J. Lü, "Characterizing and detecting inefficient image displaying issues in android apps," in *Proc. IEEE 26th Int. Conf. Softw. Anal., Evol. Reeng*, 2019, pp. 355–365.
- [80] Y. Lin, C. Radoi, and D. Dig, "Retrofitting concurrency for android applications through refactoring," in *Proc. 22nd ACM SIGSOFT Int. Symp. Found. Softw. Eng.*, 2014, pp. 341–352.
- [81] S. Okur, D. L. Hartveld, D. Dig, and A. v. Deursen, "A study and toolkit for asynchronous programming in c#," in *Proc. 36th Int. Conf. Softw. Eng.*, 2014, pp. 1117–1127.
- [82] Y. Lin, S. Okur, and D. Dig, "Study and refactoring of android asynchronous programming (T)," in *Proc. 30th IEEE/ACM Int. Conf. Automated Softw. Eng.*, 2015, pp. 224–235.
- [83] R. Feng, G. Meng, X. Xie, T. Su, Y. Liu, and S.-W. Lin, "Learning performance optimization from code changes for android apps," in *Proc. IEEE Int. Conf. Softw. Testing, Verification Validation Workshops*, 2019, pp. 285–290.
- [84] B. D. Higgins, J. Flinn, T. J. Giuli, B. Noble, C. Peplin, and D. Watson, "Informed mobile prefetching," in *Proc. 10th Int. Conf. Mobile Syst., Appl., Serv.*, 2012, pp. 155–168.
- [85] Y. Zhao, M. S. Laser, Y. Lyu, and N. Medvidovic, "Leveraging program analysis to reduce user-perceived latency in mobile applications," in *Proc. 40th Int. Conf. Softw. Eng.*, 2018, pp. 176–186.
- [86] B. Choi, J. Kim, D. Cho, S. Kim, and D. Han, "Appx: An automated app acceleration framework for low latency mobile app," in *Proc. 14th Int. Conf. Emerg. Netw. Experiments Technol.*, 2018, pp. 27–40.
- [87] I. Malavolta, F. Nocera, P. Lago, and M. Mongiello, "Navigation-aware and personalized prefetching of network requests in android apps," in *Proc. IEEE/ACM 41st Int. Conf. Softw. Eng.: New Ideas Emerg. Results*, 2019, pp. 17–20.
- [88] L. Batyuk, A.-D. Schmidt, H.-G. Schmidt, A. Camtepe, and S. Albayrak, "Developing and benchmarking native linux applications on android," in *Int. Conf. Mobile Wireless Middleware, Operating Syst., Appl.*, 2009, pp. 381–392.
- [89] S. Lee and J. W. Jeon, "Evaluating performance of android platform using native C for embedded systems," in *Proc. Int. Conf. Control, Automat., Syst.*, 2010, pp. 1160–1163.
- [90] J. K. Lee and J. Y. Lee, "Android programming techniques for improving performance," in *Proc. 3rd Int. Conf. Awareness Sci. Technol.*, 2011, pp. 386–389.
- [91] C.-M. Lin, J.-H. Lin, C.-R. Dow, and C.-M. Wen, "Benchmark Dalvik and native code for android system," in *Proc. 2nd Int. Conf. Innov. Bio-Inspired Comput. Appl.*, 2011, pp. 320–323.
- [92] C. Wang, M. Cintra, and Y. Wu, "Acceldroid: Co-designed acceleration of android bytecode," in *Proc. IEEE/ACM Int. Symp. Code Gener. Optim.*, 2013, pp. 1–10.
- [93] K.-T. T. Cheng, X. Yang, and Y.-C. Wang, "Performance optimization of vision apps on mobile application processor," in *Proc. 20th Int. Conf. Syst., Signals Image Process.*, 2013, pp. 187–191.
- [94] S. Thongkaew, T. Isshiki, D. Li, and H. Kunieda, "Dalvik bytecode acceleration using fetch/decode hardware extension," *J. Inf. Process.*, vol. 23, no. 2, pp. 118–130, 2015.
- [95] B. Mao, J. Zhou, S. Wu, H. Jiang, X. Chen, and W. Yang, "Improving flash memory performance and reliability for smartphones with i/o deduplication," *IEEE Trans. Comput.-Aided Des. Integr. Circuits Syst.*, vol. 38, no. 6, pp. 1017–1027, Jun. 2019.
- [96] H. Kim and D. Shin, "Optimizing storage performance of android smartphone," in *Proc. 7th Int. Conf. Ubiquitous Inf. Manage. Commun.*, 2013, pp. 1–7.
- [97] R. Vallée-Rai, P. Co, E. Gagnon, L. Hendren, P. Lam, and V. Sundaresan, "Soot: A java bytecode optimization framework," in *Proc. CASCON 1st Decade High Impact Papers*, 2010, pp. 214–224.
- [98] K. Nagata, S. Yamaguchi, and H. Ogawa, "A power saving method with consideration of performance in android terminals," in *Proc. 9th Int. Conf. Ubiquitous Intell. Comput. 9th Int. Conf. Autonomic Trusted Comput.*, 2012, pp. 578–585.
- [99] S. Georgiou, M. Kechagia, P. Louridas, and D. Spinellis, "What are your programming language's energy-delay implications?," in *Proc. IEEE 15th Int. Conf. Mining Softw. Repositories*, 2018, pp. 303–313.
- [100] K. Nagata, Y. Nakamura, S. Nomura, and S. Yamaguchi, "Measuring and improving application launching performance on android devices," in *Proc. 1st Int. Symp. Comput. Netw.*, 2013, pp. 636–638.
- [101] T. Yan, D. Chu, D. Ganesan, A. Kansal, and J. Liu, "Fast app launching for mobile devices using predictive user context," in *Proc. 10th Int. Conf. Mobile Syst., Appl., Serv.*, 2012, pp. 113–126.
- [102] K. Nagata and S. Yamaguchi, "An android application launch analyzing system," in *Proc. 8th Int. Conf. Comput. Technol. Inf. Manage.*, 2012, pp. 76–81.
- [103] S.-H. Kim, J. Jeong, J.-S. Kim, and S. Maeng, "SmartLMK: A memory reclamation scheme for improving user-perceived app launch time," *ACM Trans. Embedded Comput. Syst.*, vol. 15, no. 3, pp. 1–25, 2016.
- [104] Y. Joo, J. Ryu, S. Park, and K. G. Shin, "FAST: Quick application launch on solid-state drives," in *Proc. 9th UNENIX Conf. File Storage Technol.*, 2011, pp. 259–272.
- [105] Z. Zhao, M. Zhou, and X. Shen, "SatScore: Uncovering and avoiding a principled pitfall in responsiveness measurements of app launches," in *Proc. ACM Int. Joint Conf. Pervasive Ubiquitous Comput.*, 2014, pp. 21–32.
- [106] A. Parate, M. Böhmer, D. Chu, D. Ganesan, and B. M. Marlin, "Practical prediction and prefetch for faster access to applications on mobile phones," in *Proc. ACM Int. Joint Conf. Pervasive Ubiquitous Comput.*, 2013, pp. 275–284.
- [107] L.-Y. Tang, P.-C. Hsiu, J.-L. Huang, and M.-S. Chen, "iLauncher: An intelligent launcher for mobile apps based on individual usage patterns," in *Proc. 28th Annu. ACM Symp. Appl. Comput.*, 2013, pp. 505–512.
- [108] P. Baumann and S. Santini, "Every byte counts: Selective prefetching for mobile applications," in *Proc. ACM Interactive, Mobile, Wearable Ubiquitous Technol.*, 2017, pp. 1–29.
- [109] R. Prodduturi and D. B. Phatak, "Effective handling of low memory scenarios in android using logs," *Dept. Comput. Sci. Eng., Indian Inst. Technol., Mumbai, India, Tech. Rep. MTP Rep.*, 2013.
- [110] K. Baik and J. Huh, "Balanced memory management for smartphones based on adaptive background app management," in *Proc. 18th IEEE Int. Symp. Consum. Electron.*, 2014, pp. 1–2.
- [111] K. Vimal and A. Trivedi, "A memory management scheme for enhancing performance of applications on android," in *Proc. IEEE Recent Adv. Intell. Comput. Syst.*, 2015, pp. 162–166.
- [112] A. Singh, A. V. Agrawal, and A. Kanukotla, "A method to improve application launch performance in android devices," in *Proc. Int. Conf. Internet Things Appl.*, 2016, pp. 112–115.
- [113] C. Li, J. Bao, and H. Wang, "Optimizing low memory killers for mobile devices using reinforcement learning," in *Proc. 13th Int. Wireless Commun. Mobile Comput. Conf.*, 2017, pp. 2169–2174.
- [114] M. Böhmer, B. Hecht, J. Schöning, A. Krüger, and G. Bauer, "Falling asleep with angry birds, facebook and kindle: A large scale study on mobile application usage," in *Proc. 13th Int. Conf. Human Comput. Interact. Mobile Devices Serv.*, 2011, pp. 47–56.
- [115] R. Baeza-Yates, D. Jiang, F. Silvestri, and B. Harrison, "Predicting the next app that you are going to use," in *Proc. 8th ACM Int. Conf. Web Search Data Mining*, 2015, pp. 285–294.
- [116] C. Shin, J.-H. Hong, and A. K. Dey, "Understanding and prediction of mobile application usage for smart phones," in *Proc. ACM Conf. Ubiquitous Comput.*, 2012, pp. 173–182.

- [117] X. Zou, W. Zhang, S. Li, and G. Pan, "Prophet: What app you wish to use next," in *Proc. ACM Conf. Pervasive Ubiquitous Comput. Adjunct Pub.*, 2013, pp. 167–170.
- [118] J. Sylve, A. Case, L. Marziale, and G. G. Richard, "Acquisition and analysis of volatile memory from android devices," *Digit. Invest.*, vol. 8, no. 3–4, pp. 175–184, 2012.
- [119] J.-M. Kim and J.-S. Kim, "Androbench: Benchmarking the storage performance of android-based mobile devices," in *Front. Comput. Educ.*, 2012, pp. 667–674.
- [120] S. Jeong, K. Lee, J. Hwang, S. Lee, and Y. Won, "Androstep: Android storage performance analysis tool," in *Proc. Softw. Eng. Workshop*, 2013, pp. 327–340.
- [121] J. Park and B. Choi, "Automated memory leakage detection in android based systems," *Int. J. Control Automat.*, vol. 5, no. 2, pp. 35–42, 2012.
- [122] H. Shahriar, S. North, and E. Mawangi, "Testing of memory leak in android applications," in *Proc. IEEE 15th Int. Symp. High-Assurance Syst. Eng.*, 2014, pp. 176–183.
- [123] G. Santhanakrishnan, C. Cargile, and A. Olmsted, "Memory leak detection in android applications based on code patterns," in *Proc. Int. Conf. Inf. Soc.*, 2016, pp. 133–134.
- [124] D. Amalfitano, V. Riccio, P. Tramontana, and A. R. Fasolino, "Do memories haunt you? An automated black box testing approach for detecting memory leaks in android apps," *IEEE Access*, vol. 8, pp. 12 217–12 231, Jan. 2020.
- [125] T. Gerlitz, I. Kalkov, J. F. Schommer, D. Franke, and S. Kowalewski, "Non-blocking garbage collection for real-time android," in *Proc. 11th Int. Workshop Java Technol. Real-Time Embedded Syst.*, 2013, pp. 108–117.
- [126] G. Lim, C. Min, and Y. I. Eom, "Enhancing application performance by memory partitioning in android platforms," in *Proc. IEEE Int. Conf. Consum. Electron.*, 2013, pp. 649–650.
- [127] R. Mori, S. Yamaguchi, and M. Oguchi, "Memory consumption saving by optimization of promotion condition of generational GC in android," in *Proc. IEEE 6th Global Conf. Consum. Electron.*, 2017, pp. 1–2.
- [128] S.-h. Kim, J. Jeong, and J. Lee, "Efficient memory deduplication for mobile smart devices," in *Proc. IEEE Int. Conf. Consum. Electron.*, 2014, pp. 25–26.
- [129] D. Kim, E. Lee, S. Ahn, and H. Bahn, "Improving the storage performance of smartphones through journaling in non-volatile memory," *IEEE Trans. Consum. Electron.*, vol. 59, no. 3, pp. 556–561, Aug. 2013.
- [130] S. Jeong, K. Lee, S. Lee, S. Son, and Y. Won, "I/O stack optimization for smartphones," in *Proc. USENIX Annu. Techn. Conf.*, 2013, pp. 309–320.
- [131] K. Zhong *et al.*, "Building high-performance smartphones via non-volatile memory: The swap approach," in *Proc. Int. Conf. Embedded Softw.*, 2014, pp. 1–10.
- [132] S.-H. Kim, S. Kwon, J.-S. Kim, and J. Jeong, "Controlling physical memory fragmentation in mobile systems," in *Proc. Int. Symp. Memory Manage.*, 2015, p. 1–14.
- [133] D. T. Nguyen, H. Zhao, G. Zhou, G. Peng, and G. Xing, "iRAM: Sensing memory needs of my smartphone," in *Proc. IEEE 12th Int. Conf. Wireless Mobile Comput., Netw. Commun.*, 2016, pp. 1–10.
- [134] D. Kim and H. Bahn, "Exploiting write-only-once characteristics of file data in smartphone buffer cache management," *Pervasive Mobile Comput.*, vol. 40, pp. 528–540, 2017.
- [135] S.-H. Kim, J. Jeong, and J.-S. Kim, "Application-aware swapping for mobile systems," *ACM Trans. Embedded Comput. Syst.*, vol. 16, no. 5s, pp. 1–19, 2017.
- [136] J. Kim and H. Bahn, "Analysis of smartphone I/O characteristics—toward efficient swap in a smartphone," *IEEE Access*, vol. 7, pp. 129 930–129 941, 2019.
- [137] A. Escobar De La Torre, and Y. Cheon, "Impacts of java language features on the memory performances of android apps," Univ. Texas El Paso, Tech. Rep. TR-UTEP-CS-17–84, 2017.
- [138] M. Jun, L. Sheng, Y. Shengtao, T. Xianping, and L. Jian, "LeakDAF: An automated tool for detecting leaked activities and fragments of android applications," in *Proc. IEEE 41st Annu. Comput. Softw. Appl. Conf.*, 2017, pp. 23–32.
- [139] A. Carroll and G. Heiser, "An analysis of power consumption in a smartphone," in *Proc. USENIX Annu. Tech. Conf.*, Boston, MA, USA, 2010, p. 21.
- [140] L. Zhang, M. S. Gordon, R. P. Dick, Z. M. Mao, P. Dinda, and L. Yang, "ADEL: An automatic detector of energy leaks for smart-phone applications," in *Proc. 8th IEEE/ACM/IFIP Int. Conf. Hardware/Softw. Codes. Syst. Synth.*, 2012, pp. 363–372.
- [141] G. Metri, W. Shi, M. Brockmeyer, and A. Agrawal, "BatteryExtender: An adaptive user-guided tool for power management of mobile devices," in *Proc. ACM Int. Joint Conf. Pervasive Ubiquitous Comput.*, 2014, pp. 33–43.
- [142] R. Morales, R. Saborido, F. Khomh, F. Chicano, and G. Antoniol, "EARMO: An energy-aware refactoring approach for mobile apps," *IEEE Trans. Softw. Eng.*, vol. 44, no. 12, pp. 1176–1206, Dec. 2018.
- [143] A. Shye, B. Scholbrock, and G. Memik, "Into the wild: Studying real user activity patterns to guide power optimizations for mobile architectures," in *Proc. 42nd Annu. IEEE/ACM Int. Symp. Microarchitecture*, 2009, pp. 168–178.
- [144] M. A. Hoque, M. Siekkinen, K. N. Khan, Y. Xiao, and S. Tarkoma, "Modeling, profiling, and debugging the energy consumption of mobile devices," *ACM Comput. Surv.*, vol. 48, no. 3, pp. 1–40, 2015.
- [145] J. Cito, J. Rubin, P. Stanley-Marbell, and M. Rinard, "Battery-aware transformations in mobile applications," in *Proc. 31st IEEE/ACM Int. Conf. Automated Softw. Eng.*, 2016, pp. 702–707.
- [146] W. Jung, C. Kang, C. Yoon, D. Kim, and H. Cha, "DevScope: A nonintrusive and online power analysis tool for smartphone hardware components," in *Proc. 8th IEEE/ACM/IFIP Int. Conf. Hardware/Softw. Codes. Syst. Synth.*, 2012, pp. 353–362.
- [147] A. Banerjee, L. K. Chong, S. Chattopadhyay, and A. Roychoudhury, "Detecting energy bugs and hotspots in mobile apps," in *Proc. 22nd ACM SIGSOFT Int. Symp. Found. Softw. Eng.*, 2014, pp. 588–598.
- [148] E. Cuervo *et al.*, "MAUI: Making smartphones last longer with code offload," in *Proc. 8th Int. Conf. Mobile Syst., Appl., Serv.*, 2010, pp. 49–62.
- [149] A. J. Pyles, Z. Ren, G. Zhou, and X. Liu, "SiFi: Exploiting VoIP silence for WiFi energy savings in smart phones," in *Proc. 13th Int. Conf. Ubiquitous Comput.*, 2011, pp. 325–334.
- [150] B. R. Bruce, J. Petke, M. Harman, and E. T. Barr, "Approximate oracles and synergy in software energy search spaces," *IEEE Trans. Softw. Eng.*, vol. 45, no. 11, pp. 1150–1169, Nov. 2019.
- [151] A. Ferrari, D. Gallucci, D. Puccinelli, and S. Giordano, "Detecting energy leaks in android app with poem," in *Proc. IEEE Int. Conf. Pervasive Comput. Commun. Workshops*, 2015, pp. 421–426.
- [152] M. A. Bokhari, B. Alexander, and M. Wagner, "In-vivo and offline optimisation of energy use in the presence of small energy signals: A case study on a popular android library," in *Proc. 15th EAI Int. Conf. Mobile Ubiquitous Syst.: Comput., Netw. Serv.*, 2018, pp. 207–215.
- [153] A. Pathak, Y. C. Hu, M. Zhang, P. Bahl, and Y.-M. Wang, "Fine-grained power modeling for smartphones using system call tracing," in *Proc. 6th Conf. Comput. Syst.*, 2011, pp. 153–168.
- [154] C. Yoon, D. Kim, W. Jung, C. Kang, and H. Cha, "AppScope: Application energy metering framework for android smartphone using kernel activity monitoring," in *Proc. USENIX Annu. Tech. Conf.*, 2012, pp. 387–400.
- [155] M. Linares-Vásquez, G. Bavota, C. E. B. Cárdenas, R. Oliveto, M. Di Penta, and D. Poshyvanyk, "Optimizing energy consumption of GUIs in android apps: A multi-objective approach," in *Proc. 10th Joint Meeting Found. Softw. Eng.*, 2015, pp. 143–154.
- [156] M. Wan, Y. Jin, D. Li, and W. G. Halfond, "Detecting display energy hotspots in android apps," in *Proc. IEEE 8th Int. Conf. Softw. Testing. Verification Validation*, 2015, pp. 1–10.
- [157] Y. Xiao, P. Savolainen, A. Karppanen, M. Siekkinen, and A. Ylä-Jääski, "Practical power modeling of data transmission over 802.11g for wireless applications," in *Proc. 1st Int. Conf. Energy-Efficient Comput. Netw.*, 2010, pp. 75–84.
- [158] S. Hao, D. Li, W. G. Halfond, and R. Govindan, "Estimating android applications' CPU energy usage via bytecode profiling," in *Proc. 1st Int. Workshop Green Sustain. Softw.*, 2012, pp. 1–7.
- [159] S. Hao, D. Li, W. G. Halfond, and R. Govindan, "Estimating mobile application energy consumption using program analysis," in *Proc. Int. Conf. Softw. Eng.*, 2013, pp. 92–101.
- [160] D. Li, S. Hao, W. G. Halfond, and R. Govindan, "Calculating source line level energy information for android applications," in *Proc. Int. Symp. Softw. Testing Anal.*, 2013, pp. 78–89.
- [161] R. Jabbarvand, J.-W. Lin, and S. Malek, "Search-based energy testing of android," in *Proc. IEEE/ACM 41st Int. Conf. Softw. Eng.*, 2019, pp. 1119–1130.



- [162] R. Jabbarvand, A. Sadeghi, H. Bagheri, and S. Malek, "Energy-aware test-suite minimization for android apps," in *Proc. 25th Int. Symp. Softw. Testing Anal.*, 2016, pp. 425–436.
- [163] R. Mittal, A. Kansal, and R. Chandra, "Empowering developers to estimate app energy consumption," in *Proc. 18th Annu. Int. Conf. Mobile Comput. Netw.*, 2012, pp. 317–328.
- [164] M. A. Bokhari, B. R. Bruce, B. Alexander, and M. Wagner, "Deep parameter optimisation on android smartphones for energy minimisation: A tale of woe and a proof-of-concept," in *Proc. Genet. Evol. Comput. Conf. Companion*, 2017, pp. 1501–1508.
- [165] M. A. Bokhari, L. Weng, M. Wagner, and B. Alexander, "Mind the gap – a distributed framework for enabling energy optimisation on modern smart-phones in the presence of noise, drift, and statistical insignificance," in *Proc. IEEE Congr. Evol. Comput.*, 2019, pp. 1330–1337.
- [166] M. A. Bokhari, B. Alexander, and M. Wagner, "Towards rigorous validation of energy optimisation experiments," 2020, *arXiv: 2004.04500*.
- [167] A. Saarinen, M. Siekkinen, Y. Xiao, J. K. Nurminen, M. Kempainen, and P. Hui, "SmartDiet: Offloading popular apps to save energy," *ACM SIGCOMM Comput. Commun. Rev.*, vol. 42, no. 4, pp. 297–298, 2012.
- [168] A. Y. Ding, B. Han, Y. Xiao, P. Hui, A. Srinivasan, M. Kojo, and S. Tarkoma, "Enabling energy-aware collaborative mobile data offloading for smartphones," in *Proc. IEEE Int. Conf. Sens., Commun. Netw.*, 2013, pp. 487–495.
- [169] A. Saarinen, M. Siekkinen, Y. Xiao, J. K. Nurminen, M. Kempainen, and P. Hui, "Can offloading save energy for popular apps?" in *Proc. 7th ACM Int. Workshop Mobility Evol. Internet Architecture*, 2012, pp. 3–10.
- [170] A. Khairy, H. H. Ammar, and R. Bahgat, "Smartphone energizer: Extending smartphone's battery life with smart offloading," in *Proc. 9th Int. Wireless Commun. Mobile Comput. Conf.*, 2013, pp. 329–336.
- [171] Y.-W. Kwon and E. Tilevich, "Reducing the energy consumption of mobile applications behind the scenes," in *Proc. IEEE Int. Conf. Softw. Maintenance*, 2013, pp. 170–179.
- [172] L. Corral, A. B. Georgiev, A. Sillitti, G. Succi, and T. Vachkov, "Analysis of offloading as an approach for energy-aware applications on android OS: A case study on image processing," in *Int. Conf. Mobile Web Inf. Syst.*, 2014, pp. 29–40.
- [173] R. Bolla, R. Khan, X. Parra, and M. Repetto, "Improving smartphones battery life by reducing energy waste of background applications," in *Proc. 8th Int. Conf. Next Gener. Mobile Apps, Serv. Technol.*, 2014, pp. 123–130.
- [174] H. Qian and D. Andresen, "Jade: Reducing energy consumption of android app," *Int. J. Netw. Distrib. Comput.*, vol. 3, no. 3, pp. 150–158, 2015.
- [175] N. Balasubramanian, A. Balasubramanian, and A. Venkataramani, "Energy consumption in mobile phones: A measurement study and implications for network applications," in *Proc. 9th ACM SIGCOMM Conf. Internet Meas.*, 2009, pp. 280–293.
- [176] X. Chen, A. Jindal, and Y. C. Hu, "How much energy can we save from prefetching ads?: Energy drain analysis of top 100 apps," in *Proc. Workshop Power-Aware Comput. Syst.*, 2013, p. 3.
- [177] P. Mohan, S. Nath, and O. Riva, "Prefetching mobile ads: Can advertising systems afford it?," in *Proc. 8th ACM Eur. Conf. Comput. Syst.*, 2013, pp. 267–280.
- [178] Y. Yang and G. Cao, "Prefetch-based energy optimization on smartphones," *IEEE Trans. Wireless Commun.*, vol. 17, no. 1, pp. 693–706, Jan. 2018.
- [179] K. Dutta and D. Vandermeer, "Caching to reduce mobile app energy consumption," *ACM Trans. Web*, vol. 12, no. 1, pp. 1–30, 2017.
- [180] A. Pathak, Y. C. Hu, and M. Zhang, "Bootstrapping energy debugging on smartphones: A first look at energy bugs in mobile devices," in *Proc. 10th ACM Workshop Hot Top. Netw.*, 2011, pp. 1–6.
- [181] A. Pathak, A. Jindal, Y. C. Hu, and S. P. Midkiff, "What is keeping my phone awake? Characterizing and detecting no-sleep energy bugs in smartphone apps," in *Proc. 10th Int. Conf. Mobile Syst., Appl., Serv.*, 2012, pp. 267–280.
- [182] Y. Liu, C. Xu, S.-C. Cheung, and J. Lü, "GreenDroid: Automated diagnosis of energy inefficiency for smartphone applications," *IEEE Trans. Softw. Eng.*, vol. 40, no. 9, pp. 911–940, Sep. 2014.
- [183] R. Jabbarvand and S. Malek, "μDroid: An energy-aware mutation testing framework for android," in *Proc. 11th Joint Meeting Found. Softw. Eng.*, 2017, pp. 208–219.
- [184] A. Pathak, Y. C. Hu, and M. Zhang, "Where is the energy spent inside my app? Fine grained energy accounting on smartphones with Eprof," in *Proc. 7th ACM Eur. Conf. Comput. Syst.*, 2012, pp. 29–42.
- [185] S. Anwer, A. Aggarwal, R. Purandare, and V. Naik, "Chiomancer: A tool for boosting android application performance," in *Proc. 1st Int. Conf. Mobile Softw. Eng. Syst.*, 2014, pp. 62–65.
- [186] F. Alam, P. R. Panda, N. Tripathi, N. Sharma, and S. Narayan, "Energy optimization in android applications through wakelock placement," in *Proc. Des., Automat. Test Eur. Conf. Exhib.*, 2014, pp. 1–4.
- [187] D. Li, A. H. Tran, and W. G. Halfond, "Making web applications more energy efficient for OLED smartphones," in *Proc. 36th Int. Conf. Softw. Eng.*, 2014, pp. 527–538.
- [188] B. R. Bruce, J. Petke, and M. Harman, "Reducing energy consumption using genetic improvement," in *Proc. Annu. Conf. Genet. Evol. Comput.*, 2015, pp. 1327–1334.
- [189] A. Banerjee and A. Roychoudhury, "Automated re-factoring of android apps to enhance energy-efficiency," in *Proc. IEEE/ACM Int. Conf. Mobile Softw. Eng. Syst.*, 2016, pp. 139–150.
- [190] L. Cruz, R. Abreu, and J.-N. Rouvignac, "Leafactor: Improving energy efficiency of android apps via automatic refactoring," in *Proc. IEEE/ACM 4th Int. Conf. Mobile Softw. Eng. Syst.*, 2017, pp. 205–206.
- [191] A. Banerjee, L. K. Chong, C. Ballabriga, and A. Roychoudhury, "EnergyPatch: Repairing resource leaks to improve energy-efficiency of android apps," *IEEE Trans. Softw. Eng.*, vol. 44, no. 5, pp. 470–490, May 2018.
- [192] L. Cruz and R. Abreu, "Performance-based guidelines for energy efficient mobile applications," in *Proc. IEEE/ACM 4th Int. Conf. Mobile Softw. Eng. Syst.*, 2017, pp. 46–57.
- [193] L. Cruz and R. Abreu, "Using automatic refactoring to improve energy efficiency of android apps," in *Proc. CIBSE XXI Ibero-Amer. Conf. Softw. Eng.*, 2018.
- [194] M. W. Kim, D. G. Yun, J. M. Lee, and S. G. Choi, "Battery life time extension method using selective data reception on smartphone," in *Proc. Int. Conf. Inf. Netw.*, 2012, pp. 468–471.
- [195] N. Ding, D. Wagner, X. Chen, A. Pathak, Y. C. Hu, and A. Rice, "Characterizing and modeling the impact of wireless signal strength on smartphone battery drain," *ACM SIGMETRICS Perform. Eval. Rev.*, vol. 41, no. 1, pp. 29–40, 2013.
- [196] M. Bokhari and M. Wagner, "Optimising energy consumption heuristically on android mobile phones," in *Proc. Genet. Evol. Comput. Conf. Companion*, 2016, pp. 1139–1140.
- [197] K. Rao, J. Wang, S. Yalamanchili, Y. Wardi, and Y. Handong, "Application-specific performance-aware energy optimization on android mobile devices," in *Proc. IEEE Int. Symp. High Perform. Comput. Architecture*, 2017, pp. 169–180.
- [198] M. Dong, Y.-S. K. Choi, and L. Zhong, "Power modeling of graphical user interfaces on oled displays," in *Proc. 46th Annu. Des. Automat. Conf.*, 2009, pp. 652–657.
- [199] B. Anand et al., "Adaptive display power management for mobile games," in *Proc. 9th Int. Conf. Mobile Syst., Appl., Serv.*, 2011, pp. 57–70.
- [200] C.-H. Lin, P.-C. Hsiu, and C.-K. Hsieh, "Dynamic backlight scaling optimization: A cloud-based energy-saving service for mobile streaming applications," *IEEE Trans. Comput.*, vol. 63, no. 2, pp. 335–348, 2012.
- [201] C.-H. Lin, C.-K. Kang, and P.-C. Hsiu, "Catch your attention: Quality-retaining power saving on mobile OLED displays," in *Proc. 51st ACM/EDAC/IEEE Des. Automat. Conf.*, 2014, pp. 1–6.
- [202] H. Chen, J. Wang, W. Chen, H. Qu, and W. Chen, "An image-space energy-saving visualization scheme for OLED displays," *Comput. Graph.*, vol. 38, pp. 61–68, 2014.
- [203] Y. Huang, M. Chen, L. Zhang, S. Xiao, J. Zhao, and Z. Wei, "Intelligent frame refresh for energy-aware display subsystems in mobile devices," in *Proc. Int. Symp. Low Power Electron. Des.*, 2014, pp. 369–374.
- [204] X. Chen, K. W. Nixon, H. Zhou, Y. Liu, and Y. Chen, "FingerShadow: An OLED power optimization based on smartphone touch interactions," in *Proc. 6th Workshop Power-Aware Comput. Syst.*, 2014.
- [205] K. W. Nixon, X. Chen, H. Zhou, Y. Liu, and Y. Chen, "Mobile GPU power consumption reduction via dynamic resolution and frame rate scaling," in *Proc. 6th Workshop Power-Aware Comput. Syst.*, 2014, p. 5.



- [206] S. He, Y. Liu, and H. Zhou, "Optimizing smartphone power consumption through dynamic resolution scaling," in *Proc. 21st Annu. Int. Conf. Mobile Comput. Netw.*, 2015, pp. 27–39.
- [207] H.-Y. Lin, P.-C. Hsiu, and T.-W. Kuo, "Shiftmask: Dynamic OLED power shifting based on visual acuity for interactive mobile applications," in *Proc. IEEE/ACM Int. Symp. Low Power Electron. Des.*, 2017, pp. 1–6.
- [208] G. Lee, S. Lee, G. Kim, Y. Choi, R. Ha, and H. Cha, "Improving energy efficiency of android devices by preventing redundant frame generation," *IEEE Trans. Mobile Comput.*, vol. 18, no. 4, pp. 871–884, Apr. 2019.
- [209] Y.-C. Chang, W.-M. Chen, P.-C. Hsiu, Y.-Y. Lin, and T.-W. Kuo, "LSIM : Ultra lightweight similarity measurement for mobile graphics applications," in *Proc. 56th Annu. Des. Automat. Conf.*, 2019, pp. 1–6.
- [210] H.-Y. Lin, C.-C. Hung, P.-C. Hsiu, and T.-W. Kuo, "Duet: An OLED & GPU co-management scheme for dynamic resolution adaptation," in *Proc. 55th ACM/ESDA/IEEE Des. Automat. Conf.*, 2018, pp. 1–6.
- [211] P. T. Bezerra *et al.*, "Dynamic frequency scaling on android platforms for energy consumption reduction," in *Proc. 8th ACM Workshop Perform. Monit. Meas. Heterogeneous Wireless Wired Netw.*, 2013, pp. 189–196.
- [212] Y.-M. Chang, P.-C. Hsiu, Y.-H. Chang, and C.-W. Chang, "A resource-driven DVFS scheme for smart handheld devices," *ACM Trans. Embedded Comput. Syst.*, vol. 13, no. 3, pp. 1–22, 2013.
- [213] P.-H. Tseng, P.-C. Hsiu, C.-C. Pan, and T.-W. Kuo, "User-centric energy-efficient scheduling on multi-core mobile devices," in *Proc. 51st Annu. Des. Automat. Conf.*, 2014, pp. 1–6.
- [214] P.-C. Hsiu, P.-H. Tseng, W.-M. Chen, C.-C. Pan, and T.-W. Kuo, "User-centric scheduling and governing on mobile devices with big. little processors," *ACM Trans. Embedded Comput. Syst.*, vol. 15, no. 1, pp. 1–22, 2016.
- [215] S. Li and S. Mishra, "Optimizing power consumption in multicore smartphones," *J. Parallel Distrib. Comput.*, vol. 95, pp. 124–137, 2016.
- [216] P. K. Muhuri, P. K. Gupta, and J. M. Mendel, "Person footprint of uncertainty-based CWW model for power optimization in handheld devices," *IEEE Trans. Fuzzy Syst.*, vol. 28, no. 3, pp. 558–568, Mar. 2020.
- [217] J. Han and S. Lee, "Performance improvement of Linux CPU scheduler using policy gradient reinforcement learning for android smartphones," *IEEE Access*, vol. 8, pp. 11 031–11 045, 2020.
- [218] J. Paek, J. Kim, and R. Govindan, "Energy-efficient rate-adaptive GPS-based positioning for smartphones," in *Proc. 8th Int. Conf. Mobile Syst., Appl., Serv.*, 2010, pp. 299–314.
- [219] Z. Zhuang, K.-H. Kim, and J. P. Singh, "Improving energy efficiency of location sensing on smartphones," in *Proc. 8th Int. Conf. Mobile Syst., Appl., Serv.*, 2010, pp. 315–330.
- [220] Y. Chon, E. Talipov, H. Shin, and H. Cha, "Mobility prediction-based smartphone energy optimization for everyday location monitoring," in *Proc. 9th ACM Conf. Embedded Netw. Sensor Syst.*, 2011, pp. 82–95.
- [221] T. O. Oshin, S. Poslad, and A. Ma, "Improving the energy-efficiency of GPS based location sensing smartphone applications," in *Proc. IEEE 11th Int. Conf. Trust, Secur. Privacy Comput. Commun.*, 2012, pp. 1698–1705.
- [222] L. Zhang, J. Liu, H. Jiang, and Y. Guan, "Senstrack: Energy-efficient location tracking with smartphone sensors," *IEEE Sensors J.*, vol. 13, no. 10, pp. 3775–3784, Oct. 2013.
- [223] M. Linares-Vázquez, G. Bavota, C. Bernal-Cárdenas, R. Oliveto, M. Di Penta, and D. Poshyvanyk, "Mining energy-greedy api usage patterns in android apps: An empirical study," in *Proc. 11th Working Conf. Mining Softw. Repositories*, 2014, pp. 2–11.
- [224] M.-R. Ra, J. Paek, A. B. Sharma, R. Govindan, M. H. Krieger, and M. J. Neely, "Energy-delay tradeoffs in smartphone applications," in *Proc. 8th Int. Conf. Mobile Syst., Appl., Serv.*, 2010, pp. 255–270.
- [225] J. K. Nurminen, "Parallel connections and their effect on the battery consumption of a mobile phone," in *Proc. 7th IEEE Consum. Commun. Netw. Conf.*, 2010, pp. 1–5.
- [226] A. J. Pyles, X. Qi, G. Zhou, M. Keally, and X. Liu, "Sapsm: Smart adaptive 802.11 PSM for smartphones," in *Proc. ACM Conf. Ubiquitous Comput.*, 2012, pp. 11–20.
- [227] K. Lee, J. Lee, Y. Yi, I. Rhee, and S. Chong, "Mobile data offloading: How much can WiFi deliver?," *IEEE/ACM Trans. Netw.*, vol. 21, no. 2, pp. 536–550, Apr. 2013.
- [228] C.-C. Cheng and P.-C. Hsiu, "Extend your journey: Introducing signal strength into location-based applications," in *Proc. IEEE INFOCOM*, 2013, pp. 2742–2750.
- [229] M. Siekkinen, M. A. Hoque, J. K. Nurminen, and M. Aalto, "Streaming over 3G and LTE: How to save smartphone energy in radio access network-friendly way," in *Proc. 5th Workshop Mobile Video*, 2013, pp. 13–18.
- [230] D. Li, Y. Lyu, J. Gui, and W. G. Halfond, "Automated energy optimization of http requests for mobile applications," in *Proc. 38th Int. Conf. Softw. Eng.* 2016, pp. 249–260.
- [231] X. Chen, A. Jindal, N. Ding, Y. C. Hu, M. Gupta, and R. Vannithamby, "Smartphone background activities in the wild: Origin, energy drain, and optimization," in *Proc. 21st Annu. Int. Conf. Mobile Comput. Netw.*, 2015, pp. 40–52.
- [232] M. Martins, J. Cappos, and R. Fonseca, "Selectively taming background android apps to improve battery lifetime," in *Proc. USENIX Annu. Tech. Conf.*, 2015, pp. 563–575.
- [233] D. T. Nguyen *et al.*, "Storage-aware smartphone energy savings," in *Proc. 2013 ACM Int. Joint Conf. Pervasive Ubiquitous Comput.*, 2013, pp. 677–686.
- [234] A. Hussein, M. Payer, A. Hosking, and C. A. Vick, "Impact of GC design on power and performance for android," in *Proc. 8th ACM Int. Syst. Storage Conf.*, 2015, pp. 1–12.
- [235] K. Zhong *et al.*, "Energy-efficient in-memory paging for smartphones," *IEEE Trans. Comput.-Aided Des. Integr. Circuits Syst.*, vol. 35, no. 10, pp. 1577–1590, Oct. 2016.
- [236] X. Chen, Y. Chen, Z. Ma, and F. C. Fernandes, "How is energy consumed in smartphone display applications?," in *Proc. 14th Workshop Mobile Comput. Syst. Appl.*, 2013, pp. 1–6.
- [237] M. Linares-Vázquez, G. Bavota, C. Bernal-Cárdenas, M. D. Penta, R. Oliveto, and D. Poshyvanyk, "Multi-objective optimization of energy consumption of GUIs in android apps," *ACM Trans. Softw. Eng. Methodol.*, vol. 27, no. 3, pp. 1–47, 2018.
- [238] D. Li, S. Hao, J. Gui, and W. G. Halfond, "An empirical study of the energy consumption of android applications," in *Proc. IEEE Int. Conf. Softw. Maintenance Evol.*, 2014, pp. 121–130.
- [239] M. A. Hoque, M. Siekkinen, and J. K. Nurminen, "Energy efficient multimedia streaming to mobile devices—A survey," *IEEE Commun. Surv. Tut.*, vol. 16, no. 1, pp. 579–597, Jan.–Mar. 2014.
- [240] S. K. Gudla, J. K. Sahoo, A. Singh, J. Bose, and N. Ahamed, "A systematic framework to optimize launch times of web apps," in *Proc. 26th Int. Conf. World Wide Web Companion*, 2017, pp. 785–786.
- [241] L. Corral, A. B. Georgiev, A. Sillitti, and G. Succi, "Can execution time describe accurately the energy consumption of mobile apps? An experiment in android," in *Proc. 3rd Int. Workshop Green Sustain. Softw.*, 2014, pp. 31–37.
- [242] M. Xia, W. He, X. Liu, and J. Liu, "Why application errors drain battery easily? A study of memory leaks in smartphone apps," in *Proc. Workshop Power-Aware Comput. Syst.*, 2013, pp. 1–5.
- [243] D. T. Nguyen, "Evaluating impact of storage on smartphone energy efficiency," in *Proc. ACM Conf. Pervasive Ubiquitous Comput. Adjunct Pub.*, ACM, 2013, pp. 319–324.
- [244] A. P. Miettinen and J. K. Nurminen, "Energy efficiency of mobile clients in cloud computing," *HotCloud*, vol. 10, no. 4–4, p. 19, 2010.
- [245] L. Cruz and R. Abreu, "On the energy footprint of mobile testing frameworks," *IEEE Trans. Softw. Eng.*, early access, Oct. 08, 2019, doi: [10.1109/TSE.2019.2946163](https://doi.org/10.1109/TSE.2019.2946163).
- [246] R. Saborido, F. Khomh, G. Antoniol, and Y.-G. Guéhéneuc, "Comprehension of ads-supported and paid android applications: Are they different?," in *Proc. IEEE/ACM 25th Int. Conf. Program Comprehension*, 2017, pp. 143–153.
- [247] K. Wiegiers and J. Beatty, *Software Requirements*. Redmond, WA, USA: Microsoft, 2013.
- [248] J. Karlsson, "Software requirements prioritizing," in *Proc. 2nd Int. Conf. Requirements Eng.*, 1996, pp. 110–116.
- [249] J. Karlsson, C. Wohlin, and B. Regnell, "An evaluation of methods for prioritizing software requirements," *Inf. Softw. Technol.*, vol. 39, no. 14–15, pp. 939–947, 1998.
- [250] M. Woodside, G. Franks, and D. C. Petriu, "The future of software performance engineering," in *Proc. Future Softw. Eng.*, 2007, pp. 171–187.

- [251] S. Schubert, D. Kostic, W. Zwaenepoel, and K. G. Shin, "Profiling software for energy consumption," in *Proc. IEEE Int. Conf. Green Comput. Commun.*, 2012, pp. 515–522.
- [252] A. Sinha and A. P. Chandrakasan, "Jouletrack: A web based tool for software energy profiling," in *Proc. 38th Annu. Des. Automat. Conf.*, 2001, pp. 220–225.
- [253] M. Stephenson *et al.*, "Flexible software profiling of gpu architectures," in *Proc. ACM/IEEE 42nd Annu. Int. Symp. Comput. Architecture*, 2015, pp. 185–197.
- [254] X. Zhao, K. Rodrigues, Y. Luo, D. Yuan, and M. Stumm, "Non-intrusive performance profiling for entire software stacks based on the flow reconstruction principle," in *Proc. 12th USENIX Symp. Operating Syst. Des. Implementation*, 2016, pp. 603–618.
- [255] G. Xu and A. Rountev, "Precise memory leak detection for java software using container profiling," in *Proc. 30th Int. Conf. Softw. Eng.*, 2008, pp. 151–160.
- [256] M. Mao and M. Humphrey, "A performance study on the VM startup time in the cloud," in *Proc. IEEE 5th Int. Conf. Cloud Comput.*, 2012, pp. 423–430.
- [257] H. Kaminaga, "Improving linux startup time using software resume (and other techniques)," in *Proc. Linux Symp.*, 2006, p. 17.
- [258] E. Van Emden and L. Moonen, "Java quality assurance by detecting code smells," in *Proc. 9th Working Conf. Reverse Eng.*, 2002, pp. 97–106.
- [259] J.-L. Baer and T.-F. Chen, "An effective on-chip preloading scheme to reduce data access penalty," in *Proc. ACM/IEEE Conf. Supercomput.*, 1991, pp. 176–186.
- [260] C. Sahin, L. Pollock, and J. Clause, "How do code refactorings affect energy usage?," in *Proc. 8th ACM/IEEE Int. Symp. Empirical Softw. Eng. Meas.*, 2014, pp. 1–10.
- [261] F. Chen, S. Jiang, and X. Zhang, "Smartsaver: Turning flash drive into a disk energy saver for mobile computers," in *Proc. Int. Symp. Low Power Electron. Des.*, 2006, pp. 412–417.
- [262] F. Chen, T. Luo, and X. Zhang, "CAFTL: A content-aware flash translation layer enhancing the lifespan of flash memory based solid state drives," in *Proc. 9th USENIX Conf. File Storage Technol.*, vol. 11, 2011, pp. 77–90.
- [263] M. Hauswirth and T. M. Chilimbi, "Low-overhead memory leak detection using adaptive statistical profiling," in *Proc. 11th Int. Conf. Architectural Support Program. Lang. Operating Syst.*, 2004, pp. 156–164.
- [264] V. N. Padmanabhan and J. C. Mogul, "Using predictive prefetching to improve world wide web latency," *ACM SIGCOMM Comput. Commun. Rev.*, vol. 26, no. 3, pp. 22–36, 1996.
- [265] R. Minelli and M. Lanza, "Software analytics for mobile applications—insights & lessons learned," in *Proc. 17th Eur. Conf. Softw. Maintenance Reeng.*, 2013, pp. 144–153.
- [266] O. J. Romero and S. A. Ajoju, "Adroitness: An android-based middleware for fast development of high-performance apps," 2019, *arXiv:1906.02061*.
- [267] M. Linares-Vasquez, C. Vendome, Q. Luo, and D. Poshvyanyk, "How developers detect and fix performance bottlenecks in android apps," in *Proc. IEEE Int. Conf. Softw. Maintenance Evol.*, 2015, pp. 352–361.
- [268] B. Johnson, Y. Song, E. Murphy-Hill, and R. Bowdidge, "Why don't software developers use static analysis tools to find bugs?," in *Proc. Int. Conf. Softw. Eng.*, 2013, pp. 672–681.
- [269] L. Li *et al.*, "Static analysis of android apps: A systematic literature review," *Inf. Softw. Technol.*, vol. 88, pp. 67–95, 2017.
- [270] M. Christakis and C. Bird, "What developers want and need from program analysis: An empirical study," in *Proc. 31st IEEE/ACM Int. Conf. Automated Softw. Eng.*, 2016, pp. 332–343.
- [271] A. Nistor, T. Jiang, and L. Tan, "Discovering, reporting, and fixing performance bugs," in *Proc. 10th Work. Conf. Mining Softw. Repositories*, 2013, pp. 237–246.
- [272] L. Cruz, R. Abreu, J. Grundy, L. Li, and X. Xia, "Do energy-oriented changes hinder maintainability?," in *Proc. IEEE Int. Conf. Softw. Maintenance Evol.*, 2019, pp. 29–40.
- [273] M. Linares-Vasquez, C. Vendome, M. Tufano, and D. Poshvyanyk, "How developers micro-optimize android apps," *J. Syst. Softw.*, vol. 130, pp. 1–23, 2017.
- [274] L. Cruz and R. Abreu, "Catalog of energy patterns for mobile applications," *Empirical Softw. Eng.*, vol. 24, no. 4, pp. 2209–2235, 2019.
- [275] S. S. Afjehei, T.-H. P. Chen, and N. Tsantalis, "iPerfDetector: Characterizing and detecting performance anti-patterns in iOS applications," *Empirical Softw. Eng.*, vol. 24, no. 6, pp. 3484–3513, 2019.
- [276] I. Dalmasso, S. K. Datta, C. Bonnet, and N. Nikaein, "Survey, comparison and evaluation of cross platform mobile application development tools," in *Proc. 9th Int. Wireless Commun. Mobile Comput. Conf.*, 2013, pp. 323–328.
- [277] M. E. Joorabchi, A. Mesbah, and P. Kruchten, "Real challenges in mobile app development," in *Proc. ACM/IEEE Int. Symp. Empirical Softw. Eng. Meas.*, 2013, pp. 15–24.
- [278] M. Mahmoudi and S. Nadi, "The android update problem: An empirical study," in *Proc. 15th Int. Conf. Mining Softw. Repositories*, 2018, pp. 220–230.
- [279] D. Han, C. Zhang, X. Fan, A. Hindle, K. Wong, and E. Stroulia, "Understanding android fragmentation with topic analysis of vendor-specific bugs," in *Proc. 19th Work. Conf. Reverse Eng.*, 2012, pp. 83–92.
- [280] H. Khalid, M. Nagappan, E. Shihab, and A. E. Hassan, "Prioritizing the devices to test your app on: A case study of android game apps," in *Proc. 22nd ACM SIGSOFT Int. Symp. Found. Softw. Eng.*, 2014, pp. 610–620.
- [281] X. Lu *et al.*, "Prada: Prioritizing android devices for apps by mining large-scale usage data," in *Proc. IEEE/ACM 38th Int. Conf. Softw. Eng.*, 2016, pp. 3–13.
- [282] J. Harty, "Google play console: Insightful development using android vitals and pre-launch reports," in *Proc. IEEE/ACM 6th Int. Conf. Mobile Softw. Eng. Syst.*, 2019, pp. 62–65.
- [283] R. Saborido, G. Beltrame, F. Khomh, E. Alba, and G. Antoniol, "Optimizing user experience in choosing android applications," in *Proc. IEEE 23rd Int. Conf. Softw. Anal., Evol., Reeng.*, 2016, vol. 1, pp. 438–448.
- [284] E. Guzman and W. Maalej, "How do users like this feature? A fine grained sentiment analysis of app reviews," in *Proc. IEEE 22nd Int. Requirements Eng. Conf.*, 2014, pp. 153–162.
- [285] B. Fu, J. Lin, L. Li, C. Faloutsos, J. Hong, and N. Sadeh, "Why people hate your app: Making sense of user feedback in a mobile app store," in *Proc. 19th ACM SIGKDD Int. Conf. Knowl. Discov. Data Mining*, 2013, pp. 1276–1284.
- [286] P. M. Vu, T. T. Nguyen, H. V. Pham, and T. T. Nguyen, "Mining user opinions in mobile app reviews: A keyword-based approach (t)," in *Proc. 30th IEEE/ACM Int. Conf. Automated Softw. Eng.*, 2015, pp. 749–759.
- [287] C. Gao, J. Zeng, M. R. Lyu, and I. King, "Online app review analysis for identifying emerging issues," in *Proc. IEEE/ACM 40th Int. Conf. Softw. Eng.*, 2018, pp. 48–58.
- [288] Y. Tian, M. Nagappan, D. Lo, and A. E. Hassan, "What are the characteristics of high-rated apps? A case study on free android applications," in *Proc. IEEE Int. Conf. Softw. Maintenance Evol.*, 2015, pp. 301–310.
- [289] C. Shepard, A. Rahmati, C. Tossell, L. Zhong, and P. Kortum, "LiveLab: Measuring wireless networks and smartphone users in the field," *ACM SIGMETRICS Perform. Eval. Rev.*, vol. 38, no. 3, pp. 15–20, 2011.
- [290] E. Noei, M. D. Syer, Y. Zou, A. E. Hassan, and I. Keivanloo, "A study of the relation of mobile device attributes with the user-perceived quality of android apps," *Empirical Softw. Eng.*, vol. 22, no. 6, pp. 3088–3116, 2017.
- [291] U. Lee *et al.*, "Hooked on smartphones: An exploratory study on smartphone overuse among college students," in *Proc. 32nd Annu. ACM Conf. Hum. Factors Comput. Syst.*, 2014, pp. 2327–2336.
- [292] A. Morrison, X. Xiong, M. Higgs, M. Bell, and M. Chalmers, "A large-scale study of iPhone App launch behaviour," in *Proc. CHI Conf. Hum. Factors Comput. Syst.*, 2018, pp. 1–13.
- [293] D. Yu, Y. Li, F. Xu, P. Zhang, and V. Kostakos, "Smartphone app usage prediction using points of interest," in *Proc. ACM Interact., Mobile, Wearable Ubiquitous Technol.*, 2018, pp. 1–21.
- [294] J. Huangfu, J. Cao, and C. Liu, "A context-aware usage prediction approach for smartphone applications," in *Proc. Asia-Pacific Serv. Comput. Conf.*, 2015, pp. 3–16.
- [295] C. Wohlin, P. Runeson, M. Höst, M. C. Ohlsson, B. Regnell, and A. Wesslén, *Experimentation in Software Engineering*. New York, NY, USA: Springer, 2012.
- [296] M. Allamanis, E. T. Barr, P. Devanbu, and C. Sutton, "A survey of machine learning for big code and naturalness," *ACM Comput. Surv.*, vol. 51, no. 4, 2018, Art. no. 81.



**Max Hort** is currently working toward the PhD degree in software engineering with University College London, under the supervision of professor Federica Sarro and professor Mark Harman. His research interests include software fairness, nonfunctional optimization of software, and search-based software engineering.



**Maria Kechagia** received the PhD degree from the Athens University of Economics and Business and the MSc degree from Imperial College London. She is currently a research fellow with University College London. Previously, she was a postdoctoral researcher with the Delft University of Technology. Her research interests include program analysis, software testing, automated program repair, and software analytics. She has been a program committee member of ASE, MSR, and ICSME, and a reviewer for *IEEE Transactions on*

*Software Engineering*, *Empirical Software Engineering*, and the *Journal of Systems and Software*.



**Federica Sarro** is a professor of software engineering with University College London. She has authored or coauthored several articles in prestigious international venues including *International Conference on Software Engineering*, *Foundations of Software Engineering*, *Transactions on Software Engineering*, and *Transactions on Software Engineering and Methodology*. Her research interests include predictive analytics for software engineering (SE), empirical SE, and search-based SE, with a focus on software effort

estimation, software sizing, software testing, and mobile app store analysis. She has received several international awards, including the FSE'19 ACM Distinguished Paper Award and the ACM SIGEVO HUMIES GECCO'16 Award. She has also been invited to be on several program, organization and steering committees, and editorial boards of well-known venues, such as *International Conference on Software Engineering*, *Foundations of Software Engineering*, *ACM Transactions on Software Engineering and Methodology*, *IEEE Transactions on Software Engineering*, and *IEEE Transactions on Evolutionary Computation*.



**Mark Harman** works full time with Facebook London as a research scientist with a team focussing on AI for scalable software engineering. He is also a part-time professor with University College London. Previously, he was the manager of the Facebook team that deployed Sapienz to test mobile apps, which grew out of Majicke, a start up cofounded by him and acquired by Facebook in 2017. In his more purely scientific work, he cofounded the field search based software engineering, and is also known for scientific research on source code analysis, software testing, app store analysis, and empirical software engineering. He was the recipient of the IEEE Harlan Mills Award and the ACM Outstanding Research Award in 2019 for this work. In addition to Facebook itself, his scientific work is also supported by the European Research Council, with an advanced fellowship grant, and has also been regularly and generously supported by the U.K. Engineering and Physical Sciences Research Council, with regular grants, a platform, and a program grant.

► **For more information on this or any other computing topic, please visit our Digital Library at [www.computer.org/csdl](http://www.computer.org/csdl).**